

國立臺灣大學電機資訊學院資訊工程學研究所



碩士論文

Department of Computer Science and Information Engineering

College of Electrical Engineering and Computer Science

National Taiwan University

Master's Thesis

基於大型語言模型之暫存器傳輸級電路生成：
代理式 IP 重用規劃與串接路由

Agentic IP-Reuse Planning and Cascade Routing
for LLM-Based RTL Generation

洪聿鍇

Yu-Kai Hung

指導教授：洪士灝 博士

Advisor: Shih-Hao Hung, Ph.D.

中華民國 115 年 7 月

July 2026

國立臺灣大學碩士學位論文

口試委員會審定書



基於大型語言模型之暫存器傳輸級電路生成：

代理式 IP 重用規劃與串接路由

Agentic IP-Reuse Planning and Cascade Routing
for LLM-Based RTL Generation

本論文係洪聿鎔君（R13922149）在國立臺灣大學資訊工程學研究所完成之碩士學位論文，於民國 115 年 7 月 30 日承下列考試委員審查通過及口試及格，特此證明

口試委員：_____

（指導教授）

_____	_____
_____	_____
_____	_____
_____	_____

所 長：_____





摘要

大型語言模型（LLM）已能生成語法正確的程式碼，然而在真實硬體專案中，暫存器傳輸級（RTL）設計工作多以整合與矽智財（IP）重用為主，而非從零撰寫演算法模組。本論文以 RealBench 基準測試中源自三個真實開源 IP 專案（AES 加密核心、SD 卡控制器、E203 RISC-V 處理器核心）的 60 個模組級任務為評測對象，研究如何以系統工程手段——規劃、接地（grounding）、驗證、修復與路由——提升固定開源權重模型的 RTL 生成能力。

本論文提出兩項主要貢獻。其一為兩階段代理式 IP 重用流水線：規劃代理透過 IP 重用工具（目錄檢索、候選 IP 檢視與評分、介面相容性檢查）針對每個任務的真實 IP 目錄進行規劃，並以目錄注入、封閉詞彙表對齊、完整性檢查的三階段檢查產生可稽核的結構化重用計畫，再由生成器在 Verilator 驗證與修復迴圈下產生 Verilog。在 T 為零的對照實驗中，流水線通過率達 25.0%（語法 61.7%），為同一模型直接提示（8.3% / 25.0%）的 3.0 倍，且增益完全集中於 IP 重用密集的 E203 家族。其二為兩層串接路由器，依任務性質在完整流水線與低成本直接生成之間逐任務選擇流程：在單變量消融中，解題數維持基準臂的 21/60 不變，總運算時間卻由 213 小時降至 130 小時，每解一題的 token 成本為所有實驗中最低，不及暴力混合策略的三分之一。

本論文同時以對照消融精確量化各輔助機制：語意修復快取單調提升語法率



(59.4%→66.9%)，並使通過樣本更集中於可解任務；兩種修復階段皆呈「速修或永不修」分布，精簡的修復預算（語法約 6 次、功能不超過 4 次）即可保留幾乎全部的修復成功案例，同時大幅節省運算；以黃金模組檢驗持續零分任務，發現並向上游回報一項測試平台競爭條件（已獲基準採納修復），亦確認其餘零分為真實難度；同一檢測反向檢驗通過樣本時，更查獲一項提示端洩漏——直接生成提示中出現基準參考模型名稱，使部分樣本僅靠實例化黃金模型即可通過——本文所有數據均已剔除此類樣本。整體結果顯示：下一步的突破點不在語法、規劃或模型規模，而在功能正確性——需要比失配計數更豐富的回饋訊號。


關鍵字：大型語言模型、暫存器傳輸級設計、IP 重用、代理式規劃、模型路由、硬體描述語言



Abstract

Large language models generate syntactically plausible code, but real-world RTL engineering is dominated by integration and IP reuse rather than freestanding algorithm writing. This thesis studies how far system engineering — planning, grounding, verification, repair, and routing around fixed open-weight models — can push LLM-based Verilog generation on RealBench, a benchmark of 60 module-level tasks drawn from three real open-source IP projects (an AES cipher core, an SD-card controller, and the E203 RISC-V processor core) and scored by a strict warnings-as-fatal compile-and-simulate harness.

The thesis makes two primary contributions. First, a two-stage agentic IP-reuse pipeline: a planning agent works each task through a suite of IP-reuse tools — catalog search, candidate inspection and scoring, and interface-compatibility checking — and produces a structured, auditable reuse plan grounded by construction against a per-task catalog of the modules the environment actually provides (catalog injection, closed-



vocabulary alignment, and a completeness gate), followed by a generator operating under a Verilator verify-and-repair loop that mirrors the scoring environment exactly. At temperature zero the pipeline passes 25.0% of tasks (61.7% syntax) against the same model's direct-prompt 8.3% (25.0%), a $3.0\times$ improvement concentrated entirely in the reuse-rich processor-core family. Second, a two-tier cascade router that matches task character to generation flow, sending integration-shaped tasks to the pipeline and self-contained algorithmic tasks to a cheap direct flow: in its controlled ablation it holds coverage at its base arm's 21 of 60 solved tasks while reducing total compute from 213 to 130 hours, delivering the lowest token cost per solved task of any measured arm — less than a third of brute-forcing both flows.

Controlled ablations quantify the supporting machinery with equal care: a semantic repair cache lifts syntax monotonically (59.4%→66.9%) while concentrating passing samples inside solvable tasks; both repair phases converge within their first attempts, so compact budgets (≈ 6 syntax, ≤ 4 functional) retain essentially all rescues at a fraction of the compute; and a golden-model audit of persistently failing tasks uncovered a testbench race condition — since adopted and fixed upstream by the benchmark — while confirming the remaining zeros as genuine difficulty. Turned on the passes as well, the same audit caught a prompt-side leak of reference-model names that had let directly-prompted

samples pass by wrapping the benchmark's own golden model; every reported number excludes these samples. The clearest opportunity at the end is neither syntax, nor planning, nor model scale, but functional correctness, which will reward a richer feedback signal than a mismatch count.

Keywords: Large Language Models, RTL Generation, IP Reuse, Agentic Planning, LLM Routing, Verilog







Contents

	Page
Verification Letter from the Oral Examination Committee	i
摘要	iii
Abstract	v
Contents	ix
List of Figures	xiii
List of Tables	xv
Denotation	xvii
Chapter 1 Introduction	1
1.1 Motivation	1
1.2 Problem Statement	2
1.3 Contributions	2
1.4 Thesis Organization	4
Chapter 2 Background and Related Work	5
2.1 RTL Design, Verification, and IP Reuse	5
2.2 Large Language Models for Verilog Generation	6
2.3 Verification-in-the-Loop Generation and Repair	8
2.4 Retrieval-Augmented Code Generation	9

2.5	LLM Routing and Cascades	9
2.6	Positioning	10
Chapter 3	A Two-Stage Agentic IP-Reuse Pipeline	11
3.1	Overview and Environmental Contracts	12
3.2	Stage One: The Grounded Planner	14
3.3	Stage Two: The Generator	17
3.4	Repair Mechanisms	18
3.4.1	Syntax Repair Loop and the Semantic Repair Cache	18
3.4.2	Functional Repair Loop	19
3.4.3	Diagnostic-Driven Specification Slicing	20
3.5	Summary	21
Chapter 4	Cascade Routing between Direct and Pipeline Flows	23
4.1	Motivation: Two Complementary Flows	24
4.2	Task Classes and Oracle Labels	25
4.3	The Two-Tier Cascade	28
4.4	Routing Metrics	32
Chapter 5	Experimental Setup	33
5.1	The RealBench Benchmark	33
5.2	Models and Serving Infrastructure	35
5.3	Metrics and Cost Accounting	36
5.4	Run Inventory	37
Chapter 6	Results and Analysis	39
6.1	Does Planning and Reuse Beat the Model Alone?	39



6.2	The Grounding Audit: Correct, Auditable Plans	40
6.3	Sampling Structure: pass@1 to pass@10	41
6.4	Repair-Cache Ablation: Density, Not Solvability	43
6.5	Functional Repair: a Separate, Oracle-Fed Arm	46
6.6	Repair-Budget Calibration: Fix Fast or Never	46
6.7	Routing: Matched Coverage at a Fraction of the Compute	48
6.8	Failure Taxonomy: Three Tiers of Difficulty	51
6.9	When the Score Is Wrong: Benchmark and Harness Artifacts	54
Chapter 7	Discussion, Limitations, and Conclusion	59
7.1	Discussion	59
7.2	Threats to Validity	61
7.3	Future Work	63
7.4	Conclusion	64
	References	67
	Appendix A — Prompt Templates	73
A.1	Planner System Prompt (Catalog-Grounding Core)	73
A.2	Planner View vs. Generation View of a Condensed Specification . . .	74
A.3	Specifications Routed as Uncertain	76
A.4	Generation Prompt Structure	77
A.5	Functional Repair Prompt	77
	Appendix B — Per-Module Results and Reproduction	79
B.1	Per-Module Pass Counts	79
B.2	Reproduction Commands	81







List of Figures

Figure 3.1	Complete system architecture	12
Figure 3.2	Generator verify–repair loops	19
Figure 4.1	Cascade router decision flow	28
Figure 4.2	Sample flow through the cascade	31
Figure 6.1	Headline rates by task family	40
Figure 6.2	Grounding audit of reuse decisions	41
Figure 6.3	pass@k curves across runs	42
Figure 6.4	Cache ablation metrics	45
Figure 6.5	Repair-attempt distributions	47
Figure 6.6	Routing frontier	48
Figure 6.7	Router confusion vs. oracle labels	50
Figure 6.8	Outcome taxonomy across runs	51
Figure 6.9	Verilator error classes	52
Figure 6.10	Per-module pass fraction across runs	53





List of Tables

Table 2.1	Design and specification scale across Verilog benchmarks	7
Table 2.2	Positioning against related systems	10
Table 4.1	Pipeline-only vs. direct-solvable sets	26
Table 5.1	RealBench module-level task families	34
Table 5.2	Inventory of experimental runs	37
Table 5.3	Run labels to repository directories	38
Table 6.1	Headline comparison at temperature 0	39
Table 6.2	Published RealBench results vs. this thesis	44
Table 6.3	Repair-cache ablation	44
Table 6.4	Routing ablation arms	48
Table B.1	Per-module pass counts across runs	79





Denotation

LLM	Large Language Model (大型語言模型)
RTL	Register-Transfer Level (暫存器傳輸級)
HDL	Hardware Description Language (硬體描述語言)
IP	Intellectual Property core (矽智財；可重用之硬體設計區塊)
DUT	Device Under Test (受測設計；測試平台所例化之目標模組)
FSM	Finite-State Machine (有限狀態機)
CDC	Clock-Domain Crossing (跨時脈域)
RAG	Retrieval-Augmented Generation (檢索增強生成)
AES	Advanced Encryption Standard (進階加密標準)
SDC	SD-Card Controller (SD 卡控制器；本文之任務家族名稱)
E203	Nuclei HummingBirdv2 E203 開源 RISC-V 處理器核心
vLLM	高吞吐量 LLM 推論服務框架 (PagedAttention)

pass@ k	k 次取樣中至少一次通過之無偏估計通過率
syntax rate	通過基準測試編譯（無任何錯誤與未豁免警告）之樣本比例
pass rate	同時通過編譯與測試平台功能比對之樣本比例
s/f	修復預算表記：每樣本至多 s 次語法修復與 f 次功能修復
MODDUP	Verilator 錯誤類別：重複宣告環境已提供之模組
PINNOTFOUND	Verilator 錯誤類別：例化埠不存在（違反介面契約）






Chapter 1 Introduction

1.1 Motivation

Large language models have recently shown an explosive growth of capability across nearly every domain of study and research — they draft and refactor software systems, accelerate biological discovery, and increasingly act as day-to-day research assistants — and an industry has grown around turning that capability into engineering practice. Hardware design is one of the most active new directions for this wave. Hardware description languages look, superficially, like the next domain over from software: Verilog is text, compiles, and has testbenches. But the economics of hardware make correctness qualitatively harsher — an RTL bug that reaches silicon is not patched, it is respun — and the daily reality of hardware engineering differs from the coding-benchmark ideal in a way that matters more than syntax: real RTL work is dominated by integration and reuse, not by writing fresh algorithms. A working chip is assembled from verified IP blocks; the engineer’s craft lies in instantiating existing modules correctly, honoring interface contracts pin for pin, and wiring subsystems whose specifications run to hundreds of pages [11].

Most evaluations of LLM-generated Verilog sidestep exactly this reality: they pose small, self-contained modules with no environment to integrate against [15, 19]. When we instead evaluate on RealBench [10] — 60 module-level tasks carved from three real IP




projects, scored by a strict compile-and-simulate harness — a capable open-weight model prompted directly compiles on only 25.0 % of tasks and passes 8.3 %. The gap between that number and usefulness is the subject of this thesis. Our premise is that the gap is less about model intelligence than about the system around the model: whether it is told what already exists, held to the environment’s interface contracts, verified against the scorer’s own toolchain, and spared work it predictably cannot convert.

1.2 Problem Statement

Concretely, this thesis addresses the following problem. Given a natural-language specification of one Verilog module inside an existing project — with the project’s other modules provided as source files, and a hidden-logic testbench fixing the module’s exact port contract — generate RTL that compiles warning-free with the full project and passes the testbench, using only local open-weight models. Three properties make this hard. The environment imposes fatal structural contracts (re-declaring a provided module or missing a testbench port kills the sample outright). The tasks are heterogeneous — thin wrapper modules, algorithmic datapaths, and thousand-line integrators demand different generation strategies. And compute is finite and local: every planning step, repair attempt, and sample competes for the same GPU, so a system that cannot decide where its compute converts is a system that mostly wastes it.

1.3 Contributions

This thesis makes four contributions, each evaluated on the same 60-task benchmark under controlled ablations.

- 
1. **A two-stage agentic IP-reuse pipeline** (Chapter 3): a planning agent that works the task through a suite of IP-reuse tools — catalog search, candidate inspection and scoring, and interface-compatibility verification — and emits a structured reuse plan grounded against a per-task catalog of the modules the environment actually provides, followed by a generator operating under environment-faithful verification and repair. At temperature 0 the pipeline passes 25.0 % of tasks against the same model’s direct 8.3 % ($3.0\times$; syntax $2.5\times$), with the entire gain concentrated in the reuse-rich processor-core family (Section 6.1).
 2. **A cascade router between generation flows** (Chapter 4): a two-tier router — spec-level LLM pre-routing plus a plan probe for the uncertain residue — that sends integration-shaped tasks to the pipeline and self-contained algorithmic tasks to a cheap direct flow. In its controlled ablation it holds coverage at its base arm’s 21 of 60 solved tasks while reducing total compute from 213 to 130 hours and per-solved-task token cost from 2.17 M to 1.35 M — the cheapest coverage of any measured arm (Section 6.7).
 3. **Controlled measurements of the supporting machinery** (Chapter 6): a semantic repair cache that lifts syntax monotonically (59.4 %→66.9 %) while concentrating passing samples inside solvable tasks; a repair-budget calibration showing both repair phases converge within their first attempts, so compact budgets capture nearly all attainable value at a fraction of the compute; and a functional repair loop reported as its own experimental arm because its feedback signal is the scoring oracle — an accounting line this literature rarely draws.
 4. **A benchmark- and harness-artifact audit methodology** (Section 6.9): golden implementations run through the full scoring path to adjudicate persistent zeros, and

a screen of generated RTL for benchmark-gaming patterns. The audits uncovered a testbench race that scored byte-perfect S-box implementations as total failures (root-caused and fixed upstream in the benchmark), confirmed other zeros as genuine difficulty, aligned our own repair gate with the scorer — and caught a prompt-side leak that let directly-prompted samples pass by merely instantiating the benchmark’s reference models, an exclusion now applied to every number in this thesis.

Throughout, every claim is tied to a controlled comparison, and results are reported with the same care whether they move the headline metric or correct it — the audit that caught a benchmark-gaming pattern in our own direct-flow samples is, we argue, as informative as any number it revised.

1.4 Thesis Organization

Chapter 2 reviews the background and situates the thesis in four research threads. Chapters 3 and 4 present the pipeline and the router. Chapter 5 fixes the benchmark, models, metrics, and the inventory of every experimental run. Chapter 6 reports results: the headline comparison, the ablations, the routing study, a failure taxonomy, and the benchmark-artifact audit. Chapter 7 discusses implications, states threats to validity, and concludes. Appendices collect prompt templates and per-module result tables.



Chapter 2 Background and Related Work

This chapter fixes the hardware-design background the thesis assumes (Section 2.1), then surveys the four research threads it builds on: LLMs for Verilog generation (Section 2.2), verification-in-the-loop repair (Section 2.3), retrieval-augmented code generation (Section 2.4), and LLM routing and cascades (Section 2.5). Section 2.6 positions this thesis against them.

2.1 RTL Design, Verification, and IP Reuse

Register-transfer-level (RTL) design expresses digital hardware as synchronous transfers between registers through combinational logic, written in a hardware description language — here Verilog/SystemVerilog [9]. Unlike software, RTL code carries a hard structural interface: a module’s port list is a physical contract with its environment, and simulation semantics (blocking vs. nonblocking assignment, delta-cycle scheduling) can make two logically equivalent descriptions observably different to a testbench — a subtlety that returns with force in Section 6.9. Correctness is established by tools, not review: this thesis uses Verilator [28] for lint and compiled simulation and Yosys [32] for

synthesis-based structural analysis.

Industrial hardware productivity rests on IP reuse: assembling verified, pre-existing blocks rather than rewriting them, a practice codified two decades ago [11]. A working chip team spends much of its effort not on new datapaths but on integration — instantiating existing modules and wiring their interfaces correctly. This observation is the design premise of Chapter 3: if real-world RTL work is reuse-dominated, an RTL-generating agent should plan reuse explicitly rather than synthesize every line from scratch. The open-source projects our benchmark draws from — an AES cipher core and an SD-card controller from OpenCores [24] and the HummingBirdv2 E203 RISC-V processor [20] — are themselves products of this reuse culture.

2.2 Large Language Models for Verilog Generation

Benchmarks. Evaluation infrastructure for LLM-generated Verilog began with fine-tuning-era studies [25, 29] and crystallized into standard suites: VerilogEval [15], machine-checked HDLBits-derived problems with the pass@k protocol inherited from HumanEval [5], later revised with specification-to-RTL tasks [26]; and RTLLM [19], larger open-ended designs judged on syntax, functionality, and quality. These suites consist predominantly of self-contained single modules. RTL-Repo [2] moved evaluation toward multi-file repository context, and RealBench [10] — used throughout this thesis (Section 5.1) — builds tasks from complete, real IP projects in which a target module must integrate against provided dependencies and survive a strict warnings-as-fatal harness. Table 2.1 quantifies how far apart these settings are: a RealBench module task averages an order of magnitude more code lines and synthesized cells than a VerilogEval problem,

is specified by a document two orders of magnitude longer, and — alone among the suites — instantiates submodules at all. The distinction is not cosmetic: Chapter 6 shows that the phenomena that dominate RealBench (interface contracts, reuse, integration) are precisely the ones self-contained benchmarks cannot exhibit.




Table 2.1: Design and specification scale across Verilog generation benchmarks (per-design averages, reproduced from the RealBench study [10]). “Cells” = synthesized circuit cells; “Doc. lines” = specification length.

Benchmark	Designs	Code lines	Cells	Submodules	Doc. lines
Thakur et al. [29]	17	16.6	7.9	0	1.8
RTLLM v1.1 [19]	29	56.1	48.1	0.3	22.3
RTLLM v2 [19]	50	46.1	33.7	0.2	20.8
VerilogEval-human [15]	156	15.8	31.4	0	5.7
VerilogEval-machine [15]	143	13.9	6.4	0	3.2
VerilogEval v2 [26]	156	16.1	31.4	0	16.5
RealBench-module [10]	60	241.2	520	1.6	197.3
RealBench-system [10]	4	3537.3	1.1M	14.5	2900

Models and systems. One line of work improves the generator itself: VeriGen fine-tunes on harvested Verilog corpora [29]; ChipNeMo domain-adapts LLMs across chip-design tasks [14]; RTLCoder [17], CodeV [38], SiliconMind-V1 [6], and CraftRTL [16] push open models past commercial baselines with domain-specialized training pipelines — CodeV through multi-level code summarization, SiliconMind-V1 through multi-agent distillation and debug-reasoning workflows, and CraftRTL by adding correct-by-construction non-textual representations and targeted repair data. Conversational studies such as Chip-Chat [3] probed interactive design with a human in the loop. This thesis is orthogonal to all of them: we hold the base models fixed (off-the-shelf open weights, Section 5.2) and engineer the system around the model — planning, grounding, verification, repair, and routing. Better base models slot in unchanged; indeed Section 6.8 measures exactly what a $6\times$ model scale-up does and does not buy under a fixed system.

2.3 Verification-in-the-Loop Generation and Repair



Feeding tool feedback back into generation is the field’s standard answer to the unreliability of one-shot code. In software, self-repair with execution feedback is well studied and its limits documented [21]; agentic framings add explicit reasoning-action loops [36], verbal self-reflection [27], and purpose-built agent-computer interfaces [35]. In hardware, AutoChip iterates generation against compiler and simulation errors [30]; RTL-Fixer targets syntax repair with retrieval over an error-pattern database [31]; MEIC frames RTL debug as an iterated LLM workflow [33]; and VerilogCoder couples an autonomous agent with graph-based planning and an AST/waveform-tracing tool, reporting strong VerilogEval-Human results [8].

The repair machinery of Section 3.4 belongs to this family but contributes three things the prior systems do not measure. First, a semantic repair cache that replays the model’s own verified diagnostic→diff fixes across samples — related in spirit to RTLFixer’s error-pattern retrieval, but populated online from verified successes rather than curated offline, and evaluated for what it changes (sample density) and what it cannot (task solvability, Section 6.4). Second, an explicit oracle-feedback accounting: our functional repair loop consumes the scoring testbench’s mismatch signal, so it is quarantined as a separate arm — a methodological line most repair papers leave undrawn. Third, a budget calibration showing both repair phases are fix-fast-or-never (Section 6.6), turning repair-loop depth from folklore into a measured engineering parameter.

2.4 Retrieval-Augmented Code Generation



Retrieval-augmented generation grounds model output in external sources [13]; for code, retrieval over repositories supplies cross-file context that no prompt window holds natively, whether by similar-code lookup [18] or iterative repository-level retrieval interleaved with generation [37]. Our per-task IP catalog (Section 3.1) is a deliberately narrow instance: retrieval is scoped to the task’s own dependency closure — never across tasks — and its product is not only free-form context but a closed vocabulary of instantiable modules. The planner queries this catalog through its retrieval tools, the catalog digest is additionally injected into its prompt, and every reuse decision in its output is grounded against the same vocabulary post-hoc (Section 6.2) — retrieval as constraint as well as context.

2.5 LLM Routing and Cascades

A parallel literature allocates queries across models of different cost: FrugalGPT sequences models with a learned stopping rule [4]; Hybrid LLM trains a router to send easy queries to a small model [7]; RouteLLM learns routers from human preference data [22]; AutoMix self-verifies before escalating [1]. All of these route between models on general text tasks, assuming the larger model weakly dominates in quality. Our router (Chapter 4) differs in both axes. It routes between workflows over the same model — a cheap direct flow and an expensive agentic pipeline — and the expensive flow does not dominate: on self-contained algorithmic modules the pipeline is actively harmful (Section 4.1), so routing is about matching task character to flow shape, not buying quality with compute. Its

cascade structure is also domain-specific: the mid-tier probe inspects a structured artifact the pipeline produces anyway (the plan), so the escalation step costs one planner call rather than a full duplicate generation.



2.6 Positioning

Table 2.2: This thesis against the closest related systems. “Reuse planning” = explicit grounded IP-reuse decisions; “workflow routing” = per-task choice among generation flows.

System	Repair loop	Retrieval	Reuse planning	Workflow routing	Repo-realistic eval
VeriGen [29]	–	–	–	–	–
AutoChip [30]	✓	–	–	–	–
RTLFixer [31]	✓	✓	–	–	–
VerilogCoder [8]	✓	✓	–	–	–
RepoCoder [37]	–	✓	–	–	✓
RouteLLM [22]	–	–	–	model-level	–
This thesis	✓	✓	✓	✓	✓

To our knowledge, this thesis is the first system to combine catalog-grounded IP-reuse planning with per-task workflow routing for RTL generation, and the first to evaluate the combination on a repository-realistic benchmark under controlled ablations (Table 2.2). Equally deliberately, it contributes the kind of controlled diagnostic measurements the surveyed literature rarely reports: a grounding pass that makes every reuse decision in every plan auditable, a repair cache whose effect is precisely located (sample density inside solvable tasks), a budget calibration that pins down the cost-efficient repair depth, and a golden-model audit methodology that keeps the benchmark signal trustworthy (Section 6.9).




Chapter 3 A Two-Stage Agentic IP-Reuse Pipeline

This chapter presents the first contribution of the thesis: a two-stage agentic pipeline that generates RTL by first planning which existing IP blocks to reuse and how to integrate them, and only then generating Verilog under a verify-and-repair loop.

The motivation for the two-stage split is the way large language models are habitually used — and the way that habit fails on tasks of this scale. The common practice is to hand the model everything at once: the full specification, every dependency source, and the request for final code in a single giant prompt. On a RealBench task that one-shot prompt mixes hundreds of pages of behavioral description with thousands of lines of provided RTL, and the model must simultaneously decide what to build, what already exists, and how to write it — competing objectives that confuse one another and strain the context window, with the interface details that matter most buried in the middle of it. Separating the work into a planning stage (decide the reuse and integration structure over a compact, interface-centric view) and a generation stage (write one module against an already-resolved plan) gives each LLM call a single objective and a prompt shaped for exactly that objective. The separation is better on both axes we care about: accuracy, because each stage sees the right information at the right altitude, and efficiency, because

the expensive full-context reasoning happens once in the plan rather than being re-derived in every sample and repair turn.



Section 3.1 gives the dataflow and the environmental contracts the pipeline must honor. Sections 3.2 and 3.3 describe the two stages. Section 3.4 describes the three repair mechanisms layered on the generator — the syntax repair loop with a semantic repair cache, the functional repair loop, and diagnostic-driven specification slicing. Design decisions are motivated here; their empirical evaluation is deferred entirely to Chapter 6.

3.1 Overview and Environmental Contracts





Figure 3.1: Complete architecture of the generation system. The specification is parsed into sub-specs and chunk-summarized into a merged generation plan (bottom center); the generator LLM produces RTL inside a Verilator/Yosys syntax repair loop backed by the semantic repair cache, followed by testbench-driven functional repair (top); interface and behavior slices of the specification feed the two repair loops. The key-attribute extraction path on the left is the cascade router of Chapter 4, which decides per task whether generation consumes the plan or proceeds directly.

The pipeline processes one benchmark task end-to-end as follows (Figure 3.1):

- 
1. **Catalog construction.** The dependency sources that the evaluation environment provides for the task are indexed into a per-task reusable-IP catalog: one entry per provided module, carrying its name, port list, and source. The catalog is the pipeline's model of what the test environment supplies, and every downstream reuse decision is expressed in its vocabulary.
 2. **Planning.** The planner (Section 3.2) reads the specification and the catalog and emits a structured IP-reuse plan: requirements, module decomposition, a reuse decision per submodule (reuse a catalog entry vs. write new RTL), and integration notes.
 3. **Plan adaptation.** An adapter normalizes the plan, resolves each reuse decision to concrete source code, and assembles the generation prompt: the original specification, the signatures of reused modules, environment notes, and the testbench's DUT instantiation verbatim as the binding interface contract.
 4. **Generation and repair.** The generator (Section 3.3) produces the target module and iterates under an internal Verilator lint that compiles the generated code against the same file set the benchmark will use, feeding diagnostics back as repair prompts (Section 3.4).
 5. **Evaluation.** The final candidate is scored by the untouched RealBench Makefile as described in Section 5.1.

The two fatal contracts. Two properties of the evaluation environment shape nearly every design decision downstream, and both are fatal when violated. First, the environment compiles every Verilog file in the task directory together, so any provided dependency

module that the generated code re-declares produces a duplicate-module error (MODDUP): shipped dependencies must be instantiated, never re-implemented. (The single family-wide exception, `e203_extend_csr`, is not shipped as a file and must be inlined — the kind of irregularity that motivates grounding plans in the real catalog rather than in the model’s memory.) Second, the testbench instantiates the device under test with an exact port list; a generated module that misses one port fails with `PINNOTFOUND` regardless of how correct its internals are. The pipeline addresses the first contract by stripping re-declared modules and cataloging what is provided, and the second by injecting the testbench’s DUT instantiation into every generation and repair prompt as a hard interface contract.

3.2 Stage One: The Grounded Planner

The planner is a tool-calling agent: it plans by interacting with a suite of IP-reuse tools scoped to the per-task catalog, and returns a structured JSON plan with a fixed schema — requirement summary, module decomposition, per-module reuse decisions (`selected_ip` naming a catalog entry, or `new_rtl_required`), and integration, verification, and debugging plans.

The IP-reuse tool suite. Four tools carry the reuse workflow. `search_reuse_ip` queries the per-task catalog for candidate modules (token-overlap, semantic-embedding, or hybrid retrieval backends, selected by `--planner-search-mode`); `inspect_reuse_ip` returns one candidate’s behavior, interfaces, parameters, limits, and integration notes; `evaluate_ip_candidate` scores a candidate on function match, interface compatibility, configurability, and verification status; and `check_port_compatibility` parses port declarations and verifies that a pro-

posed instantiation's connections are legal. Auxiliary tools (`write_artifact`, `generate_rtl_module`, `validate_verilog`) let the agent persist plan artifacts and sanity-check RTL sketches. The intended interaction per submodule is `search` → `inspect` → `evaluate` → `decide`: the agent searches the catalog with a behavioral query, inspects the top candidates' interfaces, scores the fit, and only then commits the submodule to a reuse decision.

Reuse or new RTL. The per-module decision between `selected_ip` and `new_rtl_required` is the plan's load-bearing output, and its two error directions cost very differently. Declaring reuse of a module that does not fit (or does not exist under that name) plants a fatal seed downstream: the generator either re-implements a shipped dependency (MODDUP) or instantiates a phantom (MISSING-MODULE), both of which kill the sample at compile time. Missing a legitimate reuse is gentler but wasteful: the generator re-derives, from prose, logic that already exists in verified form — exactly the work IP reuse is meant to eliminate — and every re-derived line is a fresh opportunity for functional error. The decision rule therefore leans on evidence: a submodule is marked `selected_ip` only when a catalog entry matches both its function and its interface obligations, and `new_rtl_required` otherwise, with the grounding mechanisms below guaranteeing that whichever path produced the name, it resolves to a real catalog entry. Because the environment ships dependencies as files, reuse decisions also determine prompt composition: reused modules enter the generation prompt as interfaces to instantiate, new-RTL modules as behavioral requirements to implement.

Three mechanisms make the plan's grounding independent of any single step of the agent's trajectory — whatever mixture of tool evidence and direct reading

produced a decision, the result is checked against the catalog. Together we call them grounding-by-construction:




1. **Catalog injection.** The real catalog vocabulary — every reusable module id with a one-line digest — is injected directly into the planner’s prompt, and the system prompt mandates a closed vocabulary: `selected_ip` must be one of the listed ids.
2. **Closed-vocabulary grounding.** Every `selected_ip` in the returned plan is post-hoc resolved against the catalog: exact matches kept; near misses remapped by prefix/substring matching and then fuzzy string similarity (cutoff 0.6); irrecoverable names dropped and the submodule marked `new_rtl_required`. The plan additionally records the grounding outcome per decision, making plans auditable.
3. **Completeness gate.** A task whose catalog is non-empty but whose plan contains no reuse and no integration content is re-prompted once, and the more complete of the two attempts is kept.

Chapter 6 (Section 6.2) audits the effect: grounding eliminates hallucinated names from plans entirely, making every plan correct and auditable, while the end-to-end pass rate stays essentially level — the generator was already robust to imperfect names downstream, so the gain lands in plan quality and audibility rather than in raw yield.

Large specifications. Specifications beyond a token threshold (45,000 estimated tokens) are condensed before planning by a chunk-and-merge summarizer that renders two views from one manifest: a planner view (interfaces, reuse queries, and a brief behavioral digest) and a generation view (full requirements for to-be-generated logic; interface-only


for provided IP). Port tables and module headers are carried verbatim from the raw specification past the LLM summarizer, because interface fidelity is exactly what the two fatal contracts punish. Condensation compresses the largest specifications by roughly $5\times$ (e.g., $\sim 720\text{ k} \rightarrow \sim 146\text{ k}$ characters) while keeping every shared fact and module interface stated exactly once.



3.3 Stage Two: The Generator

The generator receives the adapted plan and produces the target module with a code-focused prompt assembled from: the (possibly condensed) specification; the source and signatures of every reused module; the environment notes (which modules are provided as files, which macros the environment defines); and the testbench DUT instantiation as the interface contract. Generated text is parsed for a Verilog code block; provided dependency modules that the model re-declared anyway are stripped before verification.

Internal verification mirrors the benchmark. The generator’s inner loop lints each candidate with Verilator against the same file set the benchmark compiles — dependencies, reference model, and testbench, with the task Makefile’s warning waivers applied — rather than linting the generated file in isolation. This design decision followed a concrete failure: an earlier isolated lint could not see testbench-side pin connections, so interface violations (PINNOTFOUND) sailed through internal checks and died only at evaluation. Mirroring the benchmark’s compile closes that gap; on the runs where both verdicts exist for identical code, the internal lint agrees with the benchmark’s syntax verdict on 60 of 60 tasks. The cost of this fidelity is that Verilator diagnostics can quote testbench or reference lines into repair prompts; we return to this leak surface in Section 7.2.



Syntax first, then function. The verification order — establish syntax correctness, then check function — is a deliberately robust flow, and it mirrors how an RTL engineer works: code that does not compile cannot be simulated, so functional evidence only exists for compiling candidates. The ordering also matches the character of the two feedback signals. Syntax diagnostics are precise, local, and actionable (a file, a line, an identifier, an error class), so they are worth iterating on first and cheaply; functional mismatch reports are coarse and global, so they are spent only on candidates that have already cleared the structural contracts. Every repair mechanism in the next section is built around this two-gate structure: candidates flow through the syntax gate into the functional gate, and no repair turn ever advances a candidate that regresses the earlier gate.

3.4 Repair Mechanisms

Three mechanisms sit on top of the generator’s inner loop (Figure 3.2). Throughout the thesis, a repair budget “ s/f ” means at most s syntax-repair and f functional-repair attempts per sample.

3.4.1 Syntax Repair Loop and the Semantic Repair Cache

When the internal lint fails, the diagnostics are fed back to the model in a repair prompt together with the current candidate, and the process repeats up to the syntax budget. The loop is enhanced by a semantic repair cache that reuses the model’s own past successes: whenever a repair subsequently passes verification, the pair (diagnostic signature $\rightarrow$ unified-diff excerpt of the fix) is stored. Signatures are Verilator diagnostic lines with file paths and identifiers stripped, gated by error code, so that structurally identical er-




Figure 3.2: The generator’s two-gate verify–repair structure. The generation plan and the specification’s interface- and behavior-topic slices feed the agentic generator; each candidate is first verified with Yosys/Verilator, whose diagnostics drive the syntax repair loop, and only then simulated against the testbench, whose mismatch reports drive the functional repair loop (off by default; Section 3.4.2). Interface-topic slices feed the syntax loop and behavior-topic slices the functional loop (Section 3.4.3).

rors on different modules collide. On a later failure, sufficiently similar cached entries are injected into the repair prompt — advisory only, never auto-applied, and never sourced from benchmark reference code. The cache has three scopes: off, task (shared only across one task’s samples), and run (shared across the whole run). Run scope maximizes hint availability but couples a task’s samples, which corrupts pass@k’s independence assumption — the trade-off is measured in Section 6.4.

3.4.2 Functional Repair Loop

The syntax loop is structurally blind to the largest failure class: candidates that compile but fail the testbench (Section 6.8). The functional repair loop attacks this class directly. After a candidate passes syntax, the testbench is run; on mismatch, the simulator’s mismatch report (“Output x has N mismatches, first at time T ”) is fed back in a prompt that frames the problem explicitly as logic repair — the interface is correct, do not change

it — and the loop keeps the best compiling candidate by lowest mismatch count, never advancing to a candidate that regresses syntax.



One methodological property makes this loop different in kind from syntax repair, and we flag it prominently: the feedback signal comes from the same testbench that scores the benchmark. A pass obtained through functional repair therefore means “specification plus oracle feedback,” not “specification alone.” For that reason the loop is off by default and reported as a separate experimental arm (Section 6.5); headline pipeline-vs-baseline comparisons never include it.

3.4.3 Diagnostic-Driven Specification Slicing

Repair prompts initially carried the full specification, which for large tasks buries the relevant sentence under tens of thousands of irrelevant tokens. Specification slicing exploits the fact that Verilator diagnostics and testbench mismatch reports name the offending identifiers exactly. The slicer extracts those identifiers, then assembles a focused slice: the verbatim interface excerpts (port tables, module headers) plus only the specification sections that mention an offending identifier, under a character budget. When a diagnostic names no identifier (e.g., a bare parse error), the slicer returns nothing and the full specification is kept, so repair is never starved of context. Interface-topic slices feed the syntax repair loop, where the errors are overwhelmingly signal-naming and port-contract issues; behavior-topic slices feed the functional repair loop, where the model must locate mis-implemented logic.

The compression this buys is large — on the largest specification the slice is a $\sim 29\times$ reduction ($\sim 704\text{ k} \rightarrow 24\text{ k}$ characters) — but the reduction rate is only meaningful when

read together with the repair success rate: a smaller prompt that dropped the fixing sentence would be a pure regression. The budget calibration of Section 6.6 shows the slice preserves the signal that matters: with sliced repair prompts, the rescues that occur land within the first few attempts of both loops, so each rescue is reached at a fraction of the full-specification token cost per attempt.

3.5 Summary

The pipeline’s design distills to three commitments. Ground every decision: the planner works over the per-task catalog through its reuse tools, and every reuse decision it emits is additionally resolved against that catalog, so plans cite only modules the environment actually provides. Verify against the real environment: the inner lint compiles exactly what the benchmark will compile, making the two fatal contracts visible during repair rather than at scoring, and the syntax-first-then-function gate ordering spends each feedback signal where it is most actionable. Reuse the model’s own verified successes: repair hints come from diffs that already fixed a structurally identical failure. What each commitment buys is quantified in Chapter 6.






Chapter 4 Cascade Routing between Direct and Pipeline Flows

Comparing the pipeline of Chapter 3 with direct generation task by task, we notice that the two flows have sharply different economics: the pipeline’s coverage subsumes the direct flow’s, but at an order of magnitude more tokens per sample — and its planning machinery, decisive on integration problems, is measurably counterproductive on self-contained algorithm and FSM design, exactly where the cheap flow matches it. Pipeline-plan generation raises accuracy substantially on integration and routing problems, where the work is arranging and wiring existing modules; direct generation writes a self-contained datapath or state machine as one coherent piece and solves the same such tasks at a tenth of the cost. This division matches the purpose of this thesis exactly: IP-reuse problems mostly need several existing modules arranged and integrated correctly, which is precisely where the plan pays, while the residual algorithmic tasks are served equally well — and far more cheaply — by letting the model generate directly. This chapter presents the second contribution of the thesis: a per-task cascade router that predicts, before spending pipeline-scale compute, which of the two flows a task should take — the full agentic pipeline, or a cheap direct single-shot generation. Section 4.1 establishes why the two flows are complementary; Section 4.2 analyzes the two solved sets and defines

the task classes and the oracle labels used to evaluate routing; Section 4.3 describes the two-tier cascade; and Section 4.4 defines the routing-specific metrics used in Section 6.7.



4.1 Motivation: Two Complementary Flows

Running both flows on all 60 tasks (the HYBRID arm of Table 5.2) motivates routing, though not in the way we first read it. With ten samples per flow, the pipeline solves 20 tasks; the direct flow solves 11 — every one of them also solved by the pipeline.¹ Coverage, then, does not argue for keeping a direct flow at all — cost does. Wrapper and integration modules pass only under the pipeline, which supplies the exact interfaces of the submodules they must instantiate and wire. Self-contained algorithmic modules pass under direct generation too, in one coherent piece, at roughly a tenth of a pipeline sample’s token cost — and on exactly those tasks the pipeline’s planning machinery is measurably counterproductive, as the next paragraph details. One flow dominates on ability; it is far from dominant on cost per solved task.

Why planning is overhead on algorithm/FSM tasks. That the pipeline converts poorly on algorithmic and FSM tasks is measurable in the HYBRID run’s own pipeline flow: on tasks whose specification centers on a state machine, compiling samples convert to passes at 16.2 %, against 43.1 % elsewhere. This is not an implementation accident; it follows from what planning does. First, decomposition fragments coherence: a state machine or an arithmetic datapath is one tightly coupled artifact, and a plan that splits it into submodule pieces forces the generator to implement fragments and glue where a single always-


¹The raw HYBRID records show the direct flow “solving” ten further tasks, but the benchmark-gaming audit of Section 6.9 traced all ten to samples that merely instantiate the benchmark’s own reference model through a prompt-side leak; they are excluded here and throughout.

block idiom — exactly the pattern the model knows best from pretraining ~~✗~~ would have expressed the whole behavior (compiling samples that follow the plan’s multi-module decomposition pass at 12.5 %, against 35.3 % for single-module outputs). Second, the plan is a paraphrase: it restates the specification’s behavior in summarized form, and for algorithmic logic the value lives in exact details (encodings, priorities, edge conditions) that summarization inevitably coarsens; the generator then optimizes toward the paraphrase rather than the source. Third, the plan’s chief asset — verbatim submodule interfaces — is worth little on a task with nothing to instantiate, so the algorithmic task pays planning’s context and structure overhead without collecting its benefit. Routing exists to give each task only the machinery it profits from.

The brute-force deployment — run both flows on every task — is the most expensive one, and it wastes most of its budget: every direct sample spent on a wrapper module converts to nothing, and every pipeline sample spent on a small self-contained module pays an order of magnitude more tokens for a task the cheap flow solves equally well. A router that sends each task to its matching flow can, in principle, hold the attainable coverage at a fraction of the cost. The routing problem is to make that prediction before the expensive flow has been paid for.

4.2 Task Classes and Oracle Labels

Before formalizing classes, it is worth looking at the two solved sets concretely rather than only through aggregate rates. In the HYBRID run, nine tasks pass only under the pipeline, and eleven pass under direct generation as well as under the pipeline (Table 4.1). The membership is telling. The pipeline-only set is wrappers, memories,



and thin glue — SRAM and ITCM/DTCM wrappers (e203_srams, e203_itcm_ram, e203_dtcn_ram), a memory controller (e203_dtcn_ctrl), the register file and CSR-control shims (e203_exu_regfile, e203_exu_alu_csrrctrl, e203_exu_alu_rglr), the IFU integrator (e203_ifu), and a mini-decoder that instantiates the real decode logic (e203_ifu_minidec): tasks whose difficulty is wiring interfaces the direct prompt cannot see. The direct-solvable set is small self-contained modules — the AES S-boxes, key expansion, and round-constant table (aes_sbox, aes_inv_sbox, aes_key_expand_128, aes_rcon), the CRC engines (sd_crc_7, sd_crc_16), a FIFO filler (sd_fifo_tx_filler), and clock/CSR/interrupt utilities (e203_clkgate, e203_clk_ctrl, e203_extend_csr, e203_irq_sync) — modules whose behavior fits in one coherent piece and affords nothing to reuse.

Table 4.1: The pipeline-only and direct-solvable sets of the HYBRID run ($60 \times (10+10)$, ref-wrap samples excluded per Section 6.9), with structural statistics of their golden implementations (median / mean / range). “Own cells” = synthesized cells of the module’s own logic after subtracting submodule instances; regular-array wrappers count as 0, modules Yosys cannot elaborate standalone are omitted. Yield counts passing samples per solved task in the set’s own flow.

	Pipeline-only (9 tasks)	Direct-solvable (11 tasks)
Members	e203_srams, e203_itcm_ram, e203_dtcn_ram, e203_dtcn_ctrl, e203_exu_regfile, e203_exu_alu_csrrctrl, e203_exu_alu_rglr, e203_ifu, e203_ifu_minidec	aes_sbox, aes_inv_sbox, aes_key_expand_128, aes_rcon, sd_crc_7, sd_crc_16, sd_fifo_tx_filler, e203_clkgate, e203_clk_ctrl, e203_extend_csr, e203_irq_sync
Code lines	114 / 145 / 71–344	93 / 121 / 34–323
Spec size (chars)	5.4k / 8.9k / 3.6k–26k	3.5k / 3.9k / 1.6k–6.8k
Own cells	0 / 24 / 0–122	20 / 201 / 0–640
Passing yield	stable (mean 7.7/10)	exploratory (mean 4.5/10)

The structural statistics in Table 4.1 still separate the two sets and make the eventual classifier concrete, though less dramatically than the raw (pre-audit) records once suggested. Pipeline-only modules synthesize to zero own cells at the median — their logic

is their submodule wiring — and carry the longer, port-table-dominated specifications. Direct-solvable modules are briefly specified, self-contained, and carry their modest logic entirely in-house (median 20 own cells, up to 640 for the S-boxes) with no dominant submodule structure. The two sets also differ in how they pass: pipeline-solved tasks pass stably (7.7 of 10 samples on average, the interfaces being supplied rather than guessed), while direct-solved tasks pass by exploring the solution space (4.5 of 10). These are exactly the observable characteristics — delegation vs. own logic, interface-dominated vs. behavior-dominated specification — that the router’s feature extractor (Section 4.3) is built to read from the specification alone.

We formalize the division as two task classes:

Class A (structural). Wrapper, integration, and memory-encapsulation modules: the module’s own logic is thin; its difficulty is wiring externally defined submodule interfaces correctly. Class A tasks want the **pipeline**.

Class B (algorithmic). Self-contained datapath or control modules — cipher rounds, CRC, instruction decode, address generation: the difficulty is the logic itself, not integration. Class B tasks want **direct** generation — not because the direct flow solves more of them, but because it solves the solvable ones at a tenth of the cost while the pipeline’s planning overhead actively hurts.

For evaluation only, each task receives an oracle label computed from its golden implementation: the reference RTL is synthesized with Yosys, submodule instances are subtracted, and the task is labeled A if its own synthesized logic falls below a 250-cell threshold (or it is a recognized regular-array wrapper), otherwise B. The label also fixes the task’s target flow (A→pipeline, B→direct). Oracle labels use golden RTL that no

generation-time component may see; they exist so that routing accuracy can be scored, and they define the ceiling arm (oracle) in the ablation of Section 6.7.



4.3 The Two-Tier Cascade





Figure 4.1: The two-tier cascade router. Tier-0 extracts the key routing attributes from the specification alone: algorithmic signals (FSM, algorithm, core logic) route to the direct flow, structural signals (wrapper, thin pass-through, delegation) route to the pipeline flow. Only ambiguous tasks generate a plan first (Tier-1); the router then re-routes by parsing the plan’s own vocabulary, and a task routed to direct after planning discards the plan (a “wasted plan”).

The router is a cascade of two increasingly expensive tests (Figure 4.1); a task exits at the first tier that is confident.

Tier-0: spec pre-route. A feature extractor reads the specification and emits seven structured judgments: whether the module delegates to submodules, is a memory encapsulation, or is thin pass-through control (all class-A signals); whether it contains a finite-state machine, clock-domain crossing, or a datapath algorithm (class-B signals); and an estimate of the state bits the module itself must hold, plus a confidence value. The production extractor is a small LLM call that returns this JSON verdict (a regex-based keyword extractor over the same features serves as the ablation baseline — Section 6.7); the decision

rule is deliberately transparent: any algorithmic signal routes to direct, else any structural signal routes to pipeline, else — or when the extractor’s confidence is below threshold — the task falls through as uncertain.



Why an LLM extractor rather than parsing. The features themselves are simple enough that keyword matching looks tempting, but specification language is semantically misleading at exactly the decision boundary. A module can sound algorithmic while being structural: e203_ifu_minidec is described as performing “preliminary decoding,” yet it is a thin wrapper that instantiates the real decoder — a keyword extractor reads “decoding” and routes it to direct, where it has no access to the decoder interface it must wire. The reverse confusion also occurs: specifications for genuinely algorithmic FIFOs and CRC engines spend most of their text on bus-facing port tables, so surface counting of interface vocabulary misreads them as structural. Deciding the class requires reading what the module does across the whole document — whether the described behavior is delegated or computed — which is a semantic judgment, not a lexical one. The whole-spec LLM read resolves both directions, and the ablation of Section 6.7 quantifies the gap to the keyword baseline.

What falls through as uncertain. The uncertain residue is not noise; it is a coherent third type of question. In the full-cascade run, exactly five tasks fell through: sd_bd (a buffer-descriptor manager), sd_data_serial_host (the SD data-line serial engine), e203_biu (the bus interface unit), e203_exu_alu_lsugu (the load/store address-generation unit sharing the ALU datapath), and e203_exu_wbck (the write-back arbiter). Their common shape: interface-arbitration modules that sit between protocols — they neither delegate to submodules (no class-A signal fires) nor implement a nameable

algorithm or FSM (no class-B signal fires), but instead mediate, arbitrate, and translate between interfaces with modest own state. For every one of the five, the extractor returned all-false features at full confidence: the specification genuinely does not commit the module to either class, and only the plan — which makes the delegation structure explicit — can. These are precisely the tasks the second tier exists for.

Tier-1: plan probe. Uncertain tasks run the planner once — the cheapest informative pipeline component — and the router probes the resulting structured plan’s own natural-language requirement summary. The plan is a good routing feature for the same reason the condensation stage renders two views (Section 3.2): the planner view is built to surface each module’s delegation and reuse structure while the generation view carries its full behavioral requirements, so by the time a plan exists, characteristics that were ambiguous in the raw specification — does this module instantiate the work or perform it? — have been made explicit in the plan’s own vocabulary of reuse decisions and requirement summaries. The probe simply reads that vocabulary back. Decisive delegation phrasing (“wraps”, “encapsulates”, “reuses existing”, or instantiation language without any own-implementation verb) keeps the task in the pipeline, plan in hand; clear self-contained-algorithm phrasing routes it to direct, discarding the plan. The asymmetry of the default is intentional: the plan is already a sunk cost, so the probe bails to direct only when the module is clearly self-contained. A plan generated and then discarded is counted as a wasted plan — the cascade’s overhead metric.

The direct flow. The direct arm is not a bare completion: it shares the generator’s repair machinery. A directly generated candidate that fails the internal lint enters the same syntax repair loop, and (in arms where functional repair is enabled) the same functional

repair loop, with the plan-dependent prompt sections simply absent. This matters for the routing economics: direct generation costs roughly 10–12 k tokens per sample against the pipeline’s 90–153 k, so every correctly-routed class-B task buys nearly an order of magnitude of compute reduction with no loss of genuine coverage.

Figure 4.2 traces how the full-cascade run’s $60 \times 10 = 600$ samples actually move through these stages, together with the volume the repair loops handle in each flow. Tier-0 routes 55 of 60 tasks from the specification alone (32 tasks to direct, 23 to the pipeline); the remaining 5 tasks’ samples pay one planner call each, after which the probe keeps 28 samples in the pipeline and releases 22 to direct. Of the resulting 342 direct and 258 pipeline samples, the syntax repair loop delivers 205 and 176 genuinely compiling candidates respectively (a further 55 direct-flow “compiles” were ref-wraps and are invalidated; Section 6.9) — the direct flow compiles at 59.9 % with most samples clean on arrival, while the pipeline flow leans on the loop harder, as befits the giant integrators routed to it — and the functional stage converts 130 samples into passes across both flows.




Figure 4.2: How the full-cascade run’s 600 samples route through the tiers and the shared repair loops (gpt-oss-20B, budgets 6/4; ref-wrap samples excluded per Section 6.9). Tier-0 decides 55 of 60 tasks from the specification; the plan probe adjudicates the 5 uncertain tasks per-sample. Both flows run the same syntax and functional repair loops; counts give each stage’s outcomes.

4.4 Routing Metrics



Section 6.7 evaluates routing arms on four axes: coverage (solved tasks and pass@k, against the pipeline’s genuine coverage as the attainable ceiling); routing accuracy (the confusion between oracle class and chosen flow, where an A→direct misroute forfeits a would-be pass and a B→pipeline misroute forfeits the cost saving); overhead (wasted plans); and efficiency (estimated tokens and summed compute per solved task). The design goal is stated once and measured throughout: hold the attainable coverage at the lowest compute, with misroutes near zero.



Chapter 5 Experimental Setup

This chapter fixes the experimental ground rules shared by every result in Chapter 6: the benchmark and its scoring semantics (Section 5.1), the models and serving infrastructure (Section 5.2), the metrics and the cost-accounting rules we commit to (Section 5.3), and a complete inventory of the experimental runs together with their configurations (Section 5.4). The inventory table doubles as the reproducibility anchor of this thesis: every number reported later traces back to exactly one row of Table 5.2.

5.1 The RealBench Benchmark

All end-to-end evaluations use RealBench [10], a Verilog generation benchmark built from three real, open-source IP projects rather than from isolated textbook exercises. Each task provides a natural-language specification of one target module, the source files of the modules it may depend on, a golden reference implementation, and a self-checking testbench. The generator must produce the target module; the surrounding project files are supplied by the evaluation environment.

The benchmark distinguishes module-level tasks, which target one module inside an otherwise-provided project, from system-level tasks, which ask for an entire subsystem. This thesis reports on the **60 module-level tasks** summarized in Table 5.1; system-level

Table 5.1: The 60 module-level RealBench tasks used in this thesis, by source project.

Family	Source project	Tasks	Character
AES	OpenCores AES-128 cipher core	6	self-contained datapath logic
SDC	OpenCores SD-card controller	14	bus-facing control logic
E203	HummingBirdv2 E203 RISC-V core	40	deep hierarchy, reuse-rich

tasks are excluded from all headline numbers because their specifications (up to ~ 704 k characters for `e203_cpu_top`) exceed what any of our serving configurations can process without aggressive condensation, and because a 60-task universe with three distinct families is already large enough to expose the phenomena this thesis studies. The three families span a useful difficulty gradient: AES modules are small and algorithmically self-contained, SDC modules implement Wishbone-facing control logic with wide interfaces, and the E203 family covers a complete RISC-V processor core whose modules instantiate one another — the setting in which IP reuse has the most to offer.

Scoring. RealBench scores with its own per-task Makefile, and its verdict is deliberately strict. A sample is counted as a syntax pass only if Verilator [28] compiles the full project — generated module, provided dependencies, reference model, and testbench together — with **no %Error and no %Warning** outside the Makefile’s explicit waiver list; warnings are fatal. A sample is a functional pass (“pass” for short) only if it also compiles and the testbench simulation reports zero output mismatches against the golden reference. Two consequences of this design recur throughout the thesis. First, our syntax numbers are conservative relative to benchmarks that fail only on errors (Section 7.2). Second, the environment imposes two non-negotiable structural contracts on generated code: a dependency that the environment ships as a file must be instantiated, never re-declared (re-declaration is a fatal MODDUP), and the testbench’s DUT instantiation fixes the exact port list (a missing port is a fatal PINNOTFOUND). Section 3.3 describes how the pipeline

internalizes both contracts.



5.2 Models and Serving Infrastructure

All generation uses open-weight models served locally through vLLM [12] behind an OpenAI-compatible endpoint, on university/NCHC GPU nodes reached through SSH port forwarding. Three models appear in this thesis:

- **gpt-oss-20B** [23] — the workhorse model. All ablations (grounding, repair cache, repair budgets, routing) run on this model, so their comparisons are model-controlled.
- **gpt-oss-120B** [23] — used to test which findings survive a $6\times$ scale increase (Sections 6.8 and 6.7).
- **Qwen3-235B** [34] — used in supplementary repair-budget runs; it does not participate in headline comparisons.

One serving behavior shaped the implementation and is documented here because it is invisible in the results yet load-bearing for reproduction: the server caps completion length rather than rejecting a request when prompt and requested tokens together exceed the 131,072-token context window, but returns HTTP 400 when the prompt alone exceeds it — the condition behind the “not generated” bucket of Section 6.8 and one motivation for the specification condensation of Section 3.2.

5.3 Metrics and Cost Accounting



Quality. For a run with n samples per task, we report the syntax rate and pass rate as per-sample means over all generated records, and pass@k per task with the standard unbiased estimator [5]

$$\text{pass@}k = \mathbb{E}_{\text{tasks}} \left[1 - \binom{n-c}{k} / \binom{n}{k} \right],$$

where c is the number of passing samples for the task, averaged over the 60 tasks. We additionally report solved tasks: the number of tasks with at least one passing sample, which equals $60 \cdot \text{pass@}n$ and is the quantity a practitioner who can afford n attempts actually cares about. Pass@k carries an independence assumption across a task’s samples; the run-scoped repair cache deliberately shares information across a task’s samples (Section 6.4), so all load-bearing pass@k values come from arms whose samples are independent, and run-scoped numbers are shown only to quantify that specific sharing effect.

Cost. We use two cost measures. Estimated tokens are computed from prompt/response character counts with a calibrated 4.38-characters-per-token heuristic and summed per run; compute time (Σ_{wall}) is the sum of per-sample wall-clock seconds, which measures total machine work independently of the harness’s ~ 80 -way request concurrency. Three accounting conventions apply throughout: (i) token totals from runs that were resumed after interruption are not cited as costs; (ii) token totals for the direct flow are reported as lower bounds; and (iii) wall-clock is compared only within a run — within-run sums and per-solved-task ratios — never across runs executed on different days on the shared server.



5.4 Run Inventory

Table 5.2 lists every experimental run this thesis reports, the configuration that distinguishes it, and the section that consumes it. Short labels (left column) are used throughout Chapter 6. All runs share the same task discovery, per-task catalog construction, and RealBench evaluation harness, so rows differ only in the flow and the flagged configuration. Unless stated otherwise, sampled runs use temperature 0.2 (the runner default) and the two-stage pipeline of Chapter 3; “6/4” style entries abbreviate the syntax/functional repair-attempt budgets of Section 3.4.

Table 5.2: Inventory of all reported runs. Every number in Chapter 6 traces to one row. Directory names are given in Table 5.3.

Label	Model	$60 \times n$	Temp.	Distinguishing configuration
<u>Baselines and headline pipeline</u>				
DIRECT-T0	20B	$\times 1$	0	spec $\rightarrow$ model, no plan, no repair
PIPE-T0	20B	$\times 1$	0	grounded planner + generator, cache off
<u>Repair-cache trio (identical except cache scope)</u>				
CACHE-OFF	20B	$\times 20$	0.2	repair cache off
CACHE-TASK	20B	$\times 20$	0.2	repair cache, task scope
CACHE-RUN	20B	$\times 20$	0.2	repair cache, run scope
<u>Model scale</u>				
PIPE-120B	120B	$\times 10$	0.2	as CACHE-RUN, $6\times$ larger model
<u>Functional repair and budgets</u>				
FUNC-2/4	20B	$\times 10$	0.2	functional repair + spec slice, budgets 2/4
FUNC-10/10	20B	$\times 10$	0.2	as above, budgets raised to 10/10
<u>Routing (Chapter 4 arms)</u>				
HYBRID	20B	$\times (10+10)$	0.2	no router: both flows on every task
ROUTER-T1	20B	$\times 10$	0.2	plan-probe tier only, budgets 6/4
ROUTER-2T	20B	$\times 10$	0.2	full cascade (Tier-0 LLM + probe), 6/4
ROUTER-2T-120B	120B	$\times 10$	0.2	as ROUTER-2T

Two rows deserve advance warning because they differ from their comparison partners in more than one variable. FUNC-10/10 changed both the repair budgets (2/4 $\rightarrow$

Table 5.3: Mapping from run labels to repository directories under runs/.

Label	Directory
DIRECT-T0	realbench_direct_model
PIPE-T0	agentic_plan_legacy_realbench_plan_hallu_fix_t0
CACHE-OFF	agentic_plan_legacy_realbench_oss20b_plan_hallu_no_rep_cache
CACHE-TASK	agentic_plan_legacy_realbench_oss20b_plan_hallu_task_rep_cache
CACHE-RUN	agentic_plan_legacy_realbench_oss20b_plan_hallu
PIPE-120B	agentic_plan_legacy_realbench_oss120b_plan_hallu_tool_call
FUNC-2/4	repair_spec_slice_oss20b_func_v2
FUNC-10/10	repair_spec_slice_oss20b_func_v2_time-err-off_10syn_10func
HYBRID	hybrid_oss20b_10syn_10func_rep_spec
ROUTER-T1	T1_router_oss20B_sync6_func4
ROUTER-2T	Full-2T_router_oss20B_sync6_func4
ROUTER-2T-120B	Full-2T_router_oss120B_sync6_func4

10/10) and a scoring waiver (TIMESCALEMOD warnings suppressed) relative to FUNC-2/4; Section 6.6 disentangles the two. And HYBRID predates the direct-flow repair loop that the router runs include, besides using the larger 10/10 budgets; Section 6.7 therefore treats ROUTER-T1 $\rightarrow$ ROUTER-2T as the only single-variable routing comparison and labels the comparison against HYBRID as directional.

One scoring rule applies to every run with a direct flow: samples whose generated RTL merely instantiates a leaked ref_* reference model are invalidated — counted as syntax and functional failures — before any rate is computed. Section 6.9 documents the leak, the detector, and the per-run impact.



Chapter 6 Results and Analysis

This chapter evaluates the pipeline (Sections 6.1 through 6.6), the router (Section 6.7), and then turns the instruments on the failures themselves (Sections 6.8 and 6.9). All runs are identified by the labels of Table 5.2; all rates were recomputed directly from each run’s per-sample records rather than transcribed from run logs, and every rate excludes ref-wrap samples — generated candidates that pass by instantiating the benchmark’s own reference model through a prompt-side leak, a benchmark-gaming pattern the audit of Section 6.9 details and quantifies.

6.1 Does Planning and Reuse Beat the Model Alone?

The cleanest head-to-head is deterministic: one greedy sample per task, identical harness, PIPE-T0 versus DIRECT-T0 (Table 6.1, Figure 6.1).

Table 6.1: Grounded pipeline vs. pure model on the 60 module tasks, temperature 0, one sample per task. Like every number in this chapter, rates exclude ref-wrap samples (Section 6.9); the exclusion costs DIRECT-T0 one syntax and one pass.

Arm	Syntax	Pass
PIPE-T0 (grounded planner + generator)	37/60 = 61.7 %	15/60 = 25.0 %
DIRECT-T0 (pure model)	15/60 = 25.0 %	5/60 = 8.3 %

The pipeline delivers $2.5\times$ the syntax rate and $3.0\times$ the pass rate of the same model prompted directly. The family breakdown locates nearly the entire effect. On AES —




Figure 6.1: Syntax and pass rates by family, pure model vs. grounded pipeline (temperature 0, one sample). The aggregate gain is carried almost entirely by the reuse-rich E203 family.

small, self-contained, nothing to reuse — the model alone compiles two thirds (66.7 %) and the pipeline’s pass gain is one task (16.7 % $\rightarrow$ 33.3 % of six). On SDC the pipeline helps modestly (pass 21.4 % $\rightarrow$ 28.6 %). On E203, the family whose modules instantiate one another, the pipeline is decisive: syntax 10.0 % $\rightarrow$ 55.0 %, pass 2.5 % $\rightarrow$ 22.5 %. The conclusion we draw — and will lean on again when routing — is that the pipeline’s value is proportional to how much legitimate IP reuse a task affords. Where there is nothing to ground, planning is overhead; where the task is an exercise in wiring real submodules, injecting their true interfaces roughly decouples the syntax rate.

6.2 The Grounding Audit: Correct, Auditable Plans

Grounding-by-construction (Section 3.2) does exactly what it was built to do. In PIPE-T0, the 60 plans contain 109 reuse decisions: 105 already name an exact catalog id, 4 near misses are remapped, and 0 hallucinations survive (Figure 6.2). At twenty samples (CACHE-RUN), the same mechanisms process 1,934 decisions: 1,834 exact, 100 remapped, 0 leaked. Before the fix, plans routinely carried invented module names; after it, every




Figure 6.2: Grounding outcomes for every reuse decision in two runs. Near-miss names are remapped onto real catalog ids; nothing hallucinated survives into the plan.

reuse decision in every plan resolves to a real, provided module — and the planner’s reuse matrix, previously empty because decisions hid under a dozen ad-hoc JSON keys, became fully populated and auditable.

The end-to-end pass rate, meanwhile, stays level: $\text{pass}@1$ moved from 19.6 % (ungrounded twin run) to 20.2 % — +0.6 percentage points, within noise. We report this measurement deliberately, because it locates the mechanism’s value precisely. The generator was already robust to imperfect names — discarding them and re-deriving reuse from the real dependency set — so grounding purchased plan correctness and audibility: every reuse decision now resolves to a real module and can be inspected, which is what the router’s plan probe (Chapter 4) and the reuse-matrix analyses build on. Plan quality was not the binding constraint on yield; Sections 6.8 and 6.9 show what is.

6.3 Sampling Structure: $\text{pass}@1$ to $\text{pass}@10$

Because most arms sample ten times per task, drawing the sampling structure from $\text{pass}@1$ to $\text{pass}@10$ puts every run on one common axis (Figure 6.3); the 60×20 arms enter through the same unbiased estimator. Read together, the curves expose three facts that





Figure 6.3: Unbiased pass@k for $k = 1-10$ across five arms on one axis. Normalizing at $k \leq 10$ lets every 60×10 run be read together with the 60×20 runs (whose curves the estimator subsamples). System design separates the curves more than model scale does.

single-run views hide. First, diversity helps every system by a similar factor and saturates fast: from $k=1$ to $k=10$ each arm gains $1.4-1.9\times$ (CACHE-OFF $18.5\% \rightarrow 29.2\%$, FUNC-10/10 $27.5\% \rightarrow 38.3\%$, HYBRID $17.7\% \rightarrow 33.1\%$, ROUTER-2T $21.7\% \rightarrow 35.0\%$), with most of each gain realized by $k=5$; extending CACHE-OFF to its full twenty samples adds only $+2.5$ pp more (31.7% at $k=20$), so samples beyond ten buy little anywhere. Second, the curves rarely cross: arms are ordered at $k=10$ almost as at $k=1$, and system differences dominate scale differences — the 20B FUNC-10/10 curve runs above the 120B pipeline curve (PIPE-120B, $24.5\% \rightarrow 35.0\%$) at every k and tops the chart. Third, the task universe splits sharply beneath every curve: on CACHE-OFF, 19 of 60 tasks are solved at least once in twenty tries and the remaining 41 never pass at any k — the sampling headroom lives entirely inside the solvable set. One decoding note carries over from the deterministic arms: greedy decoding beats the sampled per-sample mean (PIPE-T0 25.0% against $18.5-20.2\%$), the standard precision-versus-diversity trade; temperature is what buys the pass@10 headroom.



Against the published RealBench numbers. The same @1/@5 axis lets our arms be read directly against the RealBench study’s own evaluation of ten frontier and Verilog-specialized models [10] (Table 6.2). The published protocol matches ours closely — the same 60 module tasks, the same Makefile verdict, the same temperature (0.2) and unbiased estimator, with @1/@5 estimated from twenty samples against our ten — but prompts every model directly, so the comparison isolates what the system around the model is worth. The strongest published rates are func@1 13.5 % (GPT-4o-V) and func@5 17.3 % (DeepSeek-V3-671B). DIRECT-T0 — our 20B model under the same direct prompting — lands inside that frontier band (8.3 %), yet with the grounded pipeline and repair loops around it the identical weights roughly double the best published func@1 and func@5 (FUNC-10/10: 27.5 % and 35.6 %) and close to triple the best published syntax@1 (81.5 % vs. 29.3 %). The Verilog-specialized fine-tunes fare worst of all on this benchmark (CodeV-QW-7B func@1 0.3 %), reinforcing this section’s reading that system design separates RealBench performance more than model identity or scale does. Two caveats bound the comparison: the published runs predate the June-2026 AES S-box testbench fix (Section 7.2) — though excluding both S-box tasks moves none of our rates by more than 2.2 pp — and RealBench additionally reports formal@ k (equivalence checking against the golden model), which our simulation-based harness does not measure, so the table compares the two shared metrics only.

6.4 Repair-Cache Ablation: Density, Not Solvability

The cache trio (CACHE-OFF/-TASK/-RUN) differs in exactly one flag, giving the thesis’s cleanest three-arm ablation (Table 6.3, Figure 6.4).



Table 6.2: Module-level RealBench results published in the RealBench study [10] (direct prompting, temperature 0.2, @1/@5 estimated from 20 samples per task; o1-preview sampled once) against this thesis’s arms on the same 60 tasks and the same Makefile verdict (temperature 0.2, 10 samples; the -T0 arms are greedy single samples, so their @5 is undefined). Our rates exclude ref-wrap samples (Section 6.9); excluding the two post-fix-boundary S-box tasks as well would shift no rate by more than 2.2 pp.

Model / arm	Syntax (%)		Functional (%)	
	@1	@5	@1	@5
<u>Single LLMs, direct prompting; published rows from [10]</u>				
GPT-4-Turbo	27.0	36.2	9.4	13.8
o1-preview	28.3	—	13.3	—
Llama-3.1-8B	10.3	21.7	2.9	6.2
Llama-3.1-405B	6.6	15.7	2.8	7.6
DeepSeek-V3-671B	24.9	32.0	12.8	17.3
DeepSeek-R1-671B	13.4	30.4	8.0	16.2
RTLCoder-DS-6.7B	10.9	19.0	3.4	8.5
CodeV-QW-7B	3.6	7.4	0.3	1.0
GPT-4o	24.5	29.7	10.9	14.0
GPT-4o-V	29.3	38.1	13.5	16.2
gpt-oss-20B (DIRECT-T0, ours)	25.0	—	8.3	—
<u>This thesis: system around the same open weights</u>				
PIPE-T0 (20B, grounded pipeline)	61.7	—	25.0	—
FUNC-10/10 (20B, pipeline + repair)	81.5	90.7	27.5	35.6
PIPE-120B (120B, pipeline + repair)	76.3	89.5	24.5	32.8

Table 6.3: Repair-cache scope ablation (60×20 , gpt-oss-20B, identical configuration otherwise).

Cache scope	Syntax	pass@1	pass@20	mean repair attempts
off	59.4 %	18.5 %	31.7 %	1.12
task	64.3 %	19.6 %	31.7 %	1.32
run	66.9 %	20.2 %	31.7 %	1.61





Figure 6.4: The cache lifts syntax monotonically and drags pass@1 with it, but pass@20 does not move at all.

The cache does what a syntax-hint mechanism can do: verified fix hints lift the syntax rate monotonically (+4.9 pp at task scope, +7.5 pp at run scope) and pull per-sample pass up with it, at the cost of more repair turns taken. What it cannot do is change which tasks are solvable: pass@20 is identical to the decimal across all three scopes. The cache shifts the density of passing samples inside already-solvable tasks; functional bugs — the actual frontier — are immune to syntax-diff hints.

Two qualifications. First, these are three independently sampled runs, so the monotone ordering of syntax and pass together is suggestive rather than proof; individual deltas of 1–2 pp are within sampling noise. Second, run scope shares hints across a task’s twenty samples, violating pass@k’s independence assumption — run-scope pass@k is shown only to demonstrate that even with that inflation pressure the solvable set does not budge. **Task scope is the methodologically safe default**, and it is what later runs adopt. Token totals across this trio are not comparable (the task/run arms were resumed mid-run; Section 7.2); the one clean cost figure is CACHE-OFF’s 46.3 M estimated tokens for 1,200 samples.

6.5 Functional Repair: a Separate, Oracle-Fed Arm



The functional repair loop and specification slicing (Sections 3.4.2–3.4.3) are evaluated as their own arm because their feedback signal is the scoring testbench itself. Raising the budgets from 2/4 (FUNC-2/4) to 10/10 (FUNC-10/10) lifts syntax from 60.3 % to 81.5 % and pass from 20.2 % to 27.5 % — but at $2.7\times$ the cost (33.5 M $\rightarrow$ 89.8 M tokens). Attributing the +44 passing samples is sobering: +37 of them are candidates that, thanks to deeper syntax repair, now compile and then pass the testbench on their first functional check; only +7 net passes are attributable to the functional repair loop itself. (The two arms also differ in one scoring waiver — TIMESCALEMOD warnings suppressed in FUNC-10/10 — so the syntax delta slightly overstates the budget effect alone.) The clear winner inside this arm is therefore deeper syntax repair; the next section calibrates where each loop’s budget is best spent.

6.6 Repair-Budget Calibration: Fix Fast or Never

Both repair phases share one shape — fix fast or never — but with very different economics (Figure 6.5).

Syntax repair earns its budget. In FUNC-10/10, 219 samples compile without repair; repair rescues another 265 by attempt 9, with meaningful yield still arriving at attempts 4–6; past that point conversions become rare (98 samples reach attempt 10; 5 succeed). The independent ROUTER-2T run confirms the shape at budget 6/4: 153 born-clean, 228 repair-rescued with a real tail across all six attempts (73/53/43/23/22/14), the rest exhausted or




Figure 6.5: FUNC-10/10: outcomes by repair attempts consumed (log-scale counts, so the single-digit late-attempt buckets stay visible; empty buckets have no bar). Left: samples by syntax-repair attempts; nearly every bucket short of the cap ends syntax-clean, while the cap bucket (10) is 95 % failures. Right: compiling samples by functional-repair attempts; every rescue lands by attempt 7, and 324 samples burn the full budget converting nothing.

invalidated as ref-wraps. A syntax budget of about six captures nearly all attainable value.

Functional repair converges even faster. Of the compiling samples in FUNC-10/10 that enter the functional phase failing, every rescue the loop achieves — 19 of them — lands by attempt 7, and the large remainder (324 samples) runs the full depth without converting. Since each attempt is a full compile-plus-simulation (up to 300 s) plus an LLM call, a depth-ten functional budget spends 3,240 of 3,324 attempts past the point where rescues occur — which is exactly where the arm’s $2.7\times$ cost multiple lives, and exactly the compute a tighter cap recovers. ROUTER-2T confirms the same convergence at scale (13 rescues, all early, against 251 samples that run the depth out; on the direct flow the pipeline-tuned prompts convert 2 of 286 genuine samples). The calibration adopted for later runs follows directly: syntax cap ≈ 6 , functional cap ≤ 4 , plus a no-progress early exit (stop when the mismatch count stops dropping) — budgets that keep essentially all of the rescues at a fraction of the compute. Why the mismatch signal supports only early, easy fixes — it is coarse, and Section 6.9 shows it can even be misleading — is taken up

there.



6.7 Routing: Matched Coverage at a Fraction of the Compute

Table 6.4: Routing arms (gpt-oss-20B), ref-wrap samples excluded (Section 6.9). Solved = tasks with ≥ 1 passing sample. Σ wall sums per-sample wall-clock. Token totals marked + are lower bounds (direct-flow repair calls are not instrumented). HYBRID and FUNC-10/10 use budgets 10/10 and predate direct-flow repair; the single-variable comparison is ROUTER-T1 $\rightarrow$ ROUTER-2T.

Arm	Samples	pass@10	Solved/60	Σ wall (h)	Mtok/solved
HYBRID (both flows)	600+600	33.1 %	20	423	4.87 ⁺
FUNC-10/10 (pipeline only)	600	38.3 %	23	402	3.90
ROUTER-T1 (probe only)	600	35.0 %	21	213	2.17
ROUTER-2T (full cascade)	600	35.0 %	21	130	1.35⁺





Figure 6.6: Tasks solved vs. total compute. The full cascade router holds its base arm’s coverage at 39 % less compute — the cheapest solved task of any arm; the pipeline-only arm buys the most coverage at three times the router’s cost.

Table 6.4 and Figure 6.6 give the routing result. Three comparisons and one correction carry the argument.




Tier-1 alone routes conservatively. ROUTER-T1 exposes only the plan probe, and the probe — which by design defaults to keeping the sunk-cost plan — routed every task to the pipeline. The arm therefore reduces to pipeline-only at budget 6/4 (21 solved), and usefully so: it is the controlled base for the cascade, and it shows the probe needs Tier-0 in front of it to see genuinely ambiguous tasks rather than all tasks.

The cascade holds coverage while cutting compute. ROUTER-T1→ROUTER-2T changes exactly one thing: Tier-0 LLM pre-routing in front of the probe (which then handles only the 5 of 60 tasks Tier-0 deferred). Solved tasks hold at 21 and pass@10 at 35.0% — the two arms’ solved sets differ by a single low-yield task in each direction (e203_exu_disp only in the cascade, e203_exu_alu_bjp only in the base), sampling variance rather than a routing effect — while total compute falls 213 → 130 hours (−39%), because 342 of 600 samples run on the direct flow at roughly a tenth of a pipeline sample’s token cost. Per solved task, the cascade spends 1.35 M tokens against its own base arm’s 2.17 M and the brute-force hybrid’s 4.87 M: the cheapest solved task of any arm we measured.

What the raw records claimed — and why we do not. Before the benchmark-gaming audit (Section 6.9), the same records read very differently: ROUTER-2T appeared to solve 34 tasks at pass@10 56.7%, a +13-task routing win. Thirteen of those tasks were “solved” only by samples that instantiate the benchmark’s reference model — the FSM- and algorithm-heavy AES tops, SD-card engines, and E203 decode/dispatch tasks where the direct flow had seemed strongest — and HYBRID appeared to solve 30 (ten of them wrap-only). With wraps excluded, the cascade’s coverage gain over its own ablation is zero, and the direct flow contributes no task the pipeline does not also solve. The routing

result this thesis stands on is therefore the cost row of Table 6.4, not a coverage row — and we report the correction in full, because the audit that produced it is part of the contribution (Section 6.9).



Against the hybrid ceiling, directionally ahead. ROUTER-2T solves 21 tasks where HYBRID — running both flows on everything — solves 20 at $3.3\times$ the compute. This comparison is not single-variable: the hybrid predates the direct-flow repair loop (its direct samples genuinely compiled at 16.2% vs. the router run’s 59.9%) and used 10/10 budgets, so the arms differ in repair as well as routing. We therefore claim only that the cascade matches the attainable coverage at a fraction of the cost, not that it beats the hybrid deployment on coverage.




Figure 6.7: Flow allocation against oracle task classes for the full cascade, both model sizes. Misroutes concentrate in A→direct: pipeline-favoring tasks sent to the cheap flow.

Remaining headroom. Scored against the oracle labels (Figure 6.7), the 20B cascade misroutes 158 of 600 samples (26%), dominated by A→direct (110): wrapper tasks sent to the flow that cannot see their submodule interfaces — the direction that forfeits would-be passes. The probe tier is selective (5 of 60 tasks), and 22 of 279 generated plans are discarded — the cascade’s modest overhead. With gpt-oss-120B the same router solves 25

tasks at pass@1 32.7% (vs. the 20B cascade's 21 at 21.7%): scale helps genuinely, and misroutes drop to 141, differently distributed (A→direct falls to 71; B→pipeline rises to 70): the bigger model's feature extractor defers less and trusts the pipeline more. Closing the A→direct gap — e.g. by letting spec-level uncertainty trigger the probe more aggressively — is the router's clearest headroom.

6.8 Failure Taxonomy: Three Tiers of Difficulty




Figure 6.8: Per-sample outcomes. As models and repair improve, the red syntax band shrinks — and the amber compiles-but-fails-testbench band absorbs it. The functional share of compiling samples is essentially model-invariant.

Figure 6.8 decomposes every sample into four outcomes. The picture is consistent across runs: the largest failure bucket is not syntax but compiles-and-fails-the-testbench — right ports, right structure, wrong logic. Among compiling samples, that share is 59.5% in PIPE-T0, 68.9% in CACHE-OFF, and 67.9% in PIPE-120B. The 6× larger model is a strictly better Verilog syntactician — syntax 59.4% → 76.3%, first-pass compiles up, repair turns down — but it is not a better logician: the functional share of compiling samples does not move. Scale fixes syntax, not semantics; this is the strongest evidence that the functional ceiling is structural rather than a capability artifact of the 20B model.




Figure 6.9: Verilator diagnostic classes among the 442 syntax-failing samples of CACHE-OFF (a sample can carry several classes).

Reading failures per module (Figures 6.9 and 6.10) resolves the aggregate into three stable tiers:

1. **Solved: leaf, RAM, and glue modules.** `aes_key_expand_128`, the CRC and RAM blocks, clock gating, register files — small, well-specified, often pure reuse. The pipeline has essentially retired this tier.
2. **Logic-limited: datapath and control modules.** Most AES round logic and the E203 ALU/decode/commit/dispatch cluster compile with perfect syntax and fail the testbench — some at 0 passes despite dozens of structurally distinct compiling candidates. The syntax-only repair loop is blind here by construction.
3. **Scale-limited: giant integrators.** `e203_core/cpu/exu/lsu`, `sdc_controller`: hundreds of ports, dozens of submodules. Failures are interface-completeness (PINNOTFOUND concentrates exactly here), context overflow (the not-generated bucket: timeouts, empty returns, and HTTP 400s on the biggest specs), and integration logic. Notably, plan-level interface hallucination reappears only on this





Figure 6.10: Pass fraction per task (rows, grouped AES/SDC/E203) across six runs (columns). The three difficulty tiers are visible as texture: saturated rows (leaf/glue modules), intermittent rows (logic-limited datapath modules), and blank rows (giant integrators).

tier — the e203_exu plan promotes 25 internal signals to phantom top ports while dropping 10 real ones — the same interface-too-large-to-hold phenomenon at a different stage.



Each tier wants a different intervention — nothing (tier 1), functional feedback (tier 2), decomposition and contract enforcement (tier 3) — which is why aggregate pass rate alone is a blunt instrument for steering this system.

6.9 When the Score Is Wrong: Benchmark and Harness Artifacts

Treating every 0 % module as model failure would misdirect the next round of engineering, so the persistent zeros were audited individually with a golden-model methodology: run the benchmark's own reference implementation through the full scoring path, and if the golden fails (or a byte-equivalent candidate fails), the task's zeros are artifacts, not model errors. The same skepticism was then turned around and applied to the persistent passes. Four audits produced four different verdicts.

The AES S-box race (confirmed artifact, since fixed upstream). `aes_sbox` and `aes_inv_sbox` scored 0 across all runs with one suspicious signature: every sample — combinational, ROM-array, or registered variants alike — reported exactly 766 mismatches. Diffing showed generated lookup tables byte-identical to the golden. The cause is a testbench race: stimulus changes the input and samples the output in the same simulation timestep, so any correct-but-differently-scheduled implementation (e.g. a continuous assign where the reference uses blocking always `@(a))` settles one delta

cycle late at each of the ~ 768 input transitions. The trigger was isolated by incremental flag bisection to Verilator's `--coverage-line` instrumentation — the identical design passes without it — and the recorded mismatch count decomposes exactly as $766 + 3 \times (\text{wrong table bytes})$, meaning the artifact also corrupts mismatch counts as a repair signal: a one-byte-wrong design (0.4 % wrong) reports $\sim 50\%$ mismatches. A registered-output module using the same testbench template (`aes_key_expand_128`) passes 20/20, completing the case that this is a combinational-settling race. Impact: 21 byte-perfect samples scored zero in CACHE-OFF alone, and 2 of 60 tasks were unsolvable purely by artifact. The finding was reported and the benchmark's testbench was repaired upstream on June 25, 2026 (settle before compare); pre-/post-fix runs are never compared on these two tasks.

The SDC controller (audited, genuine). `sdc_controller` — 0 functional passes in 220 records — looked like a candidate for the same treatment and is the reason the golden methodology matters: its golden implementation passes its testbench perfectly (0 mismatches in 23,932 checks) under the full scoring flags, because all its outputs are registered and immune to the settling race. The zeros are honest: 87 % of SDC samples die on syntax through interface hallucination introduced at the generation stage (instantiating provided submodules with invented pins, inventing a different master-bus naming, dropping the `sd_defines` macros) even when the plan's interface is correct — and the 13 % that compile are grossly wrong ($\sim 95\%$ mismatch rates), not near misses.

The repair gate stricter than the scorer (pipeline-side finding). Auditing the 97 samples of FUNC-10/10 that exhausted all ten syntax repairs found that 42 of them (43 %) fail on a fatal warning located in a benchmark harness file (`stimulus_gen/ref`) — code

the model cannot edit. The internal repair gate (`verilator --lint-only`) treats three warning classes as fatal that the actual scorer's build does not even emit, so on four tasks every sample burned its entire repair budget "fixing" correct code and never reached the functional phase. Golden test again: three of the four goldens pass the real scoring build (pure pipeline artifact; the fix is to mirror the scorer's waivers in the gate), while the fourth (`e203_reset_ctrl`) fails even as golden — a genuine benchmark defect of the AES class. The remaining 53 % of repair-exhausted samples fail on real errors in generated code, concentrated on the giant integrators of tier 3.

The reference-wrap leak (audited, our-side; all runs re-scored). The fourth audit inverted the question: instead of asking whether zeros were honest, it asked whether the passes were. The direct flow's prompt lists the verification directory's provided helper modules so the model will not redeclare them — and that list included `ref_<task>`, the benchmark's own golden reference model. The models found the exploit: emit a top module that merely instantiates `ref_<task>`, which compiles (the scoring Makefile globs the whole verification directory) and matches the reference by construction. Detection is mechanical — a generated sample is a ref-wrap if any line instantiates a `ref_*` module — and unambiguous: every flagged sample wraps exactly its own task's reference, and none occurs in the pipeline flow, whose runner filters `ref_*` names from its prompts. The exploit's yield tracks the machinery around it: in HYBRID, whose direct flow had no repair loop, only 48 of 186 attempted wraps compiled; in ROUTER-2T, with repair, 54 of 56 passed; at 120B, all 51 attempts passed. The impact was concentrated exactly where it was most misleading: 13 of ROUTER-2T's 34 raw solved tasks and 10 of HYBRID's 30 were wrap-only, manufacturing the apparent direct-flow advantage on FSM and algorithm tasks that the routing analysis had initially credited (Section 6.7). Every number in this

thesis excludes wrap samples (counted as syntax and functional failures); the fix on the generation side is one line — filter `ref_*` from the direct prompt's provided-module list, as the pipeline runner already does — and wrap screening is now part of our scoring path.

The methodological point generalizes beyond this benchmark: in a harness this strict, persistent zeros — and flattering passes — deserve a golden-model audit before they are allowed to steer engineering effort. Of our four audits, one cleared the models (and improved the benchmark upstream), one confirmed the difficulty as genuine, one located a configuration to align in our own gate, and one caught our own prompt handing the models a shortcut and corrected a headline ablation before it could mislead. Each verdict redirected the engineering correctly.






Chapter 7 Discussion, Limitations, and Conclusion

7.1 Discussion

What actually moved the needle. Ranked by measured effect, the levers were:

(1) environment-faithful verification and repair — compiling against the real file set, injecting the testbench port contract, and condensing giant specs took the pipeline from low-single-digit to $\sim 18\text{--}25\%$ pass, the single largest step in the project's history; (2) routing — matching task character to flow shape held coverage while cutting compute by 39 % and per-solved-task tokens by roughly 40 % (Section 6.7); (3) model scale — $6\times$ parameters bought +9 pp syntax and +4 pp pass@1 with the system held fixed; (4) repair caching — +5–7 pp syntax, density only. Smaller on the headline metric but valuable elsewhere: grounding (+0.6 pp, while buying auditability and eliminating a failure class) and the functional repair loop (+7 net passes, best run at a compact budget). The general lesson we draw is unglamorous: in agentic code generation for strict environments, fidelity between the inner loop and the scorer is worth more than additional intelligence in any single stage.



Grounding as a harness property. The planner’s tool suite gives it evidence on demand — catalog search, candidate inspection, interface checks — but the design never lets plan correctness depend on any particular trajectory through those tools: the catalog vocabulary is also injected into the prompt, the output schema is constrained to that vocabulary, and every emitted reuse decision is resolved against the catalog post-hoc. The result is that plans are correct by construction regardless of how the model arrived at them, and the grounding audit (Section 6.2) can verify it decision by decision. We suspect this pattern — make correctness a property of the harness, layered over rather than delegated to the model’s tool-use behavior — transfers well beyond RTL.

Why routing pays here. The routing literature mostly buys quality with compute along one axis (small model vs. large model). Our setting is different in kind: the expensive flow’s competence subsumes the cheap flow’s, but the two differ by an order of magnitude in cost per sample, and the scaffold that makes the expensive flow strong on integration-heavy instances is measurably counterproductive on self-contained ones (Section 4.1) — so routing to the cheap path loses nothing there and saves nearly everything. We expect this shape to recur wherever an elaborate agentic scaffold coexists with a cheap direct path, because the routing signal (does this task’s difficulty live in its interfaces or in its logic?) is legible from the specification itself. The corollary from our audit is equally general: when a cheap flow suddenly looks more able than the scaffold — rather than merely cheaper — that is a claim worth auditing before it is banked (Section 6.9).

The functional ceiling. The binding constraint at the end of this work is unambiguous: ~60–69 % of everything that compiles fails its testbench, a share that neither model scale, nor caching, nor ten rounds of oracle-fed repair meaningfully dents. Verilator diagnostics

give the repair loop a precise, local, actionable signal for syntax; a mismatch count gives almost nothing — and Section 6.9 showed it can even be actively misleading. Progress on this tier will require a fundamentally better feedback channel (waveform-level localization, per-submodule benches, assertion synthesis), not more attempts on the current one.

7.2 Threats to Validity

We collect here, in one place, every caveat that qualifies the numbers in this thesis; each was also flagged inline where it bites.

1. **Single-seed runs.** Temperature-0 arms are one sample per task; each sampled run is one seed. Cross-run deltas under $\sim 2\text{--}3$ pp are compatible with sampling noise. The claims we rest weight on (headline $3.0\times$; routing -39% compute at matched coverage; the flat pass@20) are far outside that band — the routing arms' one-task set differences are inside it and are treated as noise.
2. **Cross-run configuration drift.** Only three comparisons are single-variable: the cache trio, the temp-0 head-to-head, and ROUTER-T1 $\rightarrow$ ROUTER-2T. Everything else — notably the router-vs-HYBRID comparison (repair budgets and direct-flow repair differ) and the FUNC-2/4 $\rightarrow$ FUNC-10/10 pair (a scoring waiver changed) — is labeled directional where it appears.
3. **Cost accounting.** Token totals from resumed runs undercount (resumed samples log zero tokens): the only clean full-run token cost is CACHE-OFF's 46.3 M. Direct-flow repair calls are not instrumented, so direct-flow token totals —

including ROUTER-2T’s 28.3 M — are lower bounds; the per-solved-task advantage in Table 6.4 (1.35 M vs. the base arm’s measured 2.17 M) would shrink somewhat, not reverse, under full instrumentation. Wall-clock is never compared across days (shared server, varying load).

4. **pass@k independence.** The run-scoped cache couples a task’s samples; all load-bearing pass@k values come from cache-off or task-scoped arms.
5. **Scoring strictness.** The harness fails on warnings, so our syntax rates are conservative relative to error-only benchmarks — and Section 6.9 showed our own internal gate was at times stricter than the scorer, deflating four tasks.
6. **Benchmark fix boundary.** The AES S-box testbenches were repaired upstream on June 25, 2026, after our report. Every run in this thesis predates the fix except where stated; no comparison spans the boundary on aes_sbox/aes_inv_sbox.
7. **Ref-wrap exclusion.** All rates exclude samples that instantiate a leaked ref_* reference model (Section 6.9). Detection is a syntactic screen over generated RTL; it flagged only unambiguous self-wraps and cannot, by construction, catch subtler uses of the leaked name — though the reference source was never visible to the models, bounding what a subtler exploit could copy. The direct-flow prompt retains the leak in the archived runs; the definitive control is a re-run with the filtered prompt (Section 7.3).
8. **Oracle feedback.** The functional repair loop consumes the scoring testbench’s mismatch report. Its passes mean “specification plus oracle feedback” and are quarantined in Section 6.5; no headline number includes them. Relatedly, the internal lint compiles the testbench, so Verilator diagnostics can quote testbench lines into repair


prompts; a flag exists to disable this for leak-free claims, at a syntax-rate cost we did not pay here.



9. **Benchmark scope.** All results are on one benchmark’s 60 module-level tasks and three IP projects, with open-weight models of one family for all controlled comparisons. The family-level heterogeneity we observe (AES vs. E203) is itself evidence that results will vary across codebases.

7.3 Future Work

Five directions follow directly from the measured constraints. First, functional feedback without oracle leakage: synthesize self-checking benches or assertions from the specification, so the functional repair loop gets a mismatch signal the scorer never sees — and gate it off pure-combinational modules, where Section 6.9 showed the signal can be corrupted. Second, combining the pipeline with a testbench-generation pipeline to unlock fine-grained module-decomposition IP reuse: today the plan’s decomposition stops at module granularity because only the top module has an oracle, but a companion pipeline that generates a bench per planned submodule would let the system decompose more aggressively, verify each decomposed piece independently, and reuse IP at the sub-block level — turning the planner’s decomposition from a prompt-structuring device into a verified construction process. Third, interface-contract enforcement as a hard gate: the testbench DUT instantiation is already injected as text; promoting it to a mechanical post-generation port-diff check would delete the PINNOTFOUND class that dominates giant-module syntax failures, and the internal repair gate should mirror the scorer’s warning waivers exactly. Fourth, selective decomposition for the giants: RealBench’s system-




level tasks decompose into submodules that are themselves benched module tasks, so a tree-structured generator could verify each node against a real testbench and compose upward — decomposition with per-node oracles, reserved for the scale-limited tier where it pays, and the natural first customer of the testbench-generation direction above. Fifth, hardening the router: the dominant misroute ($A \rightarrow \text{direct}$) argues for widening the uncertainty band that triggers the plan probe; a replication with the `ref_*`-filtered direct prompt is the definitive control for the ref-wrap exclusion of Section 6.9; and multi-seed replication would move the routing result from strongly suggestive to statistically settled.

7.4 Conclusion

This thesis asked how far careful system engineering — planning, grounding, verification, repair, and routing around fixed open-weight models — can push LLM-based RTL generation on a realistic, reuse-dominated benchmark. The answer is: substantially, and measurably. A two-stage pipeline that plans IP reuse against a grounded per-task catalog and generates under environment-faithful verification reaches $2.5\times$ the syntax rate and $3.0\times$ the pass rate of the same model prompted directly, with the entire gain concentrated where reuse exists to be exploited. A two-tier cascade router that matches task character to generation flow holds that coverage — matching its pipeline-only base arm task for task — at 39 % less compute and the lowest token cost per solved task of any arm measured, less than a third of the brute-force deployment's.

The thesis's second kind of contribution is measurement discipline over the same system: grounding made every plan correct and auditable; caching concentrated passing density inside solvable tasks; the repair budgets were calibrated to where the rescues

actually happen, recovering the large majority of the functional phase's compute; and a golden-model audit of persistent zeros cleared the models on two tasks, aligned our own repair gate with the scorer on four, confirmed the benchmark sound on the rest — and, turned on our own passes, caught a prompt-side leak that had inflated the routing ablation and corrected it before it could mislead. These findings redirected the engineering each time, and one was adopted upstream by the benchmark itself. The frontier, at the end, is not syntax, not plans, and not scale: it is functional correctness, and the next system to move it will be the one that finds its models a better feedback signal than a mismatch count.









References

- [1] P. Aggarwal, A. Madaan, A. Anand, S. P. Potharaju, S. Mishra, P. Zhou, A. Gupta, D. Rajagopal, K. Kappaganthu, Y. Yang, S. Upadhyay, M. Faruqui, and Mausam. AutoMix: Automatically mixing language models. In Advances in Neural Information Processing Systems 37 (NeurIPS), 2024. arXiv:2310.12963.
- [2] A. Allam and M. Shalan. RTL-Repo: A benchmark for evaluating LLMs on large-scale RTL design projects. In 2024 IEEE LLM Aided Design Workshop (LAD), 2024. arXiv:2405.17378.
- [3] J. Blocklove, S. Garg, R. Karri, and H. Pearce. Chip-chat: Challenges and opportunities in conversational hardware design. In 2023 ACM/IEEE 5th Workshop on Machine Learning for CAD (MLCAD), 2023. arXiv:2305.13243.
- [4] L. Chen, M. Zaharia, and J. Zou. FrugalGPT: How to use large language models while reducing cost and improving performance. Transactions on Machine Learning Research, 2024. First appeared as arXiv:2305.05176 (2023).
- [5] M. Chen, J. Tworek, H. Jun, Q. Yuan, H. P. de Oliveira Pinto, J. Kaplan, H. Edwards, Y. Burda, N. Joseph, G. Brockman, A. Ray, R. Puri, G. Krueger, M. Petrov, H. Khlaaf, G. Sastry, P. Mishkin, B. Chan, S. Gray, et al. Evaluating large language models trained on code, 2021.

- 
- [6] M.-C. Chen, Y.-H. Kao, P.-H. Huang, S.-C. Ho, H.-Y. Tsou, I.-T. Wu, E.-M. Huang, Y.-K. Hung, W.-P. Hsin, C. Liang, C.-H. Tu, S.-H. Hung, and H. T. Kung. SiliconMind-V1: Multi-agent distillation and debug-reasoning workflows for Verilog code generation, 2026.
- [7] D. Ding, A. Mallick, C. Wang, R. Sim, S. Mukherjee, V. Rühle, L. V. S. Lakshmanan, and A. H. Awadallah. Hybrid LLM: Cost-efficient and quality-aware query routing. In The Twelfth International Conference on Learning Representations (ICLR), 2024. arXiv:2404.14618.
- [8] C.-T. Ho, H. Ren, and B. Khailany. VerilogCoder: Autonomous Verilog coding agents with graph-based planning and abstract syntax tree (AST)-based waveform tracing tool. In Proceedings of the AAAI Conference on Artificial Intelligence, volume 39, pages 300–307, 2025. arXiv:2408.08927.
- [9] IEEE. IEEE standard for SystemVerilog—unified hardware design, specification, and verification language. IEEE Std 1800-2023, 2023.
- [10] P. Jin, D. Huang, C. Li, S. Cheng, Y. Zhao, X. Zheng, J. Zhu, S. Xing, B. Dou, R. Zhang, Z. Du, Q. Guo, and X. Hu. RealBench: Benchmarking Verilog generation models with real-world IP designs, 2025.
- [11] M. Keating and P. Bricaud. Reuse Methodology Manual for System-on-a-Chip Designs. Kluwer Academic Publishers, Boston, MA, 3rd edition, 2002.
- [12] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica. Efficient memory management for large language model serving with PagedAttention. In Proceedings of the 29th ACM Symposium on Operating Systems Principles (SOSP). ACM, 2023.

- 
- [13] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. tau Yih, T. Rocktäschel, S. Riedel, and D. Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In Advances in Neural Information Processing Systems 33 (NeurIPS), 2020. arXiv:2005.11401.
- [14] M. Liu, T.-D. Ene, R. Kirby, C. Cheng, N. Pinckney, R. Liang, J. Alben, H. Anand, S. Banerjee, I. Bayraktaroglu, et al. ChipNeMo: Domain-adapted LLMs for chip design, 2023.
- [15] M. Liu, N. Pinckney, B. Khailany, and H. Ren. VerilogEval: Evaluating large language models for Verilog code generation. In 2023 IEEE/ACM International Conference on Computer Aided Design (ICCAD), 2023. arXiv:2309.07544.
- [16] M. Liu, Y.-D. Tsai, W. Zhou, and H. Ren. CraftRTL: High-quality synthetic data generation for Verilog code models with correct-by-construction non-textual representations and targeted code repair. In The Thirteenth International Conference on Learning Representations (ICLR), 2025. arXiv:2409.12993.
- [17] S. Liu, W. Fang, Y. Lu, Q. Zhang, H. Zhang, and Z. Xie. RTLCoder: Outperforming GPT-3.5 in design RTL generation with our open-source dataset and lightweight solution. In 2024 IEEE LLM Aided Design Workshop (LAD), 2024. arXiv:2312.08617.
- [18] S. Lu, N. Duan, H. Han, D. Guo, S. won Hwang, and A. Svyatkovskiy. ReACC: A retrieval-augmented code completion framework. In Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pages 6227–6240. Association for Computational Linguistics, 2022.
- [19] Y. Lu, S. Liu, Q. Zhang, and Z. Xie. RTLLM: An open-source benchmark for design


RTL generation with large language model. In Proceedings of the 29th Asia and South Pacific Design Automation Conference (ASP-DAC), 2024.



- [20] Nuclei System Technology. Hummingbirdv2 E203 ultra-low power RISC-V processor core. https://github.com/riscv-mcu/e203_hbirdv2, 2026. Open-source project. Accessed July 2026.
- [21] T. X. Olausson, J. P. Inala, C. Wang, J. Gao, and A. Solar-Lezama. Is self-repair a silver bullet for code generation? In The Twelfth International Conference on Learning Representations (ICLR), 2024. arXiv:2306.09896.
- [22] I. Ong, A. Almahairi, V. Wu, W.-L. Chiang, T. Wu, J. E. Gonzalez, M. W. Kadous, and I. Stoica. RouteLLM: Learning to route LLMs from preference data. In The Thirteenth International Conference on Learning Representations (ICLR), 2025. arXiv:2406.18665.
- [23] OpenAI. gpt-oss-120b & gpt-oss-20b model card, 2025.
- [24] OpenCores. OpenCores: Open-source IP cores repository. <https://opencores.org>, 2026. Accessed July 2026.
- [25] H. Pearce, B. Tan, and R. Karri. DAVE: Deriving automatically Verilog from English. In Proceedings of the 2020 ACM/IEEE Workshop on Machine Learning for CAD (MLCAD), pages 27–32, 2020.
- [26] N. Pinckney, C. Batten, M. Liu, H. Ren, and B. Khailany. Revisiting VerilogEval: A year of improvements in large-language models for hardware code generation. ACM Transactions on Design Automation of Electronic Systems, 30(6):91:1–91:20, 2025. arXiv:2408.11053.

- 
- [27] N. Shinn, F. Cassano, A. Gopinath, K. Narasimhan, and S. Yao. Reflexion: Language agents with verbal reinforcement learning. In Advances in Neural Information Processing Systems 36 (NeurIPS), 2023. arXiv:2303.11366.
- [28] W. Snyder. Verilator: Open-source SystemVerilog simulator and lint system. <https://verilator.org>, 2026. Accessed July 2026.
- [29] S. Thakur, B. Ahmad, H. Pearce, B. Tan, B. Dolan-Gavitt, R. Karri, and S. Garg. VeriGen: A large language model for Verilog code generation. ACM Transactions on Design Automation of Electronic Systems, 29(3):46:1–46:31, 2024.
- [30] S. Thakur, J. Blocklove, H. Pearce, B. Tan, S. Garg, and R. Karri. AutoChip: Automating HDL generation using LLM feedback, 2023.
- [31] Y.-D. Tsai, M. Liu, and H. Ren. RTLFixer: Automatically fixing RTL syntax errors with large language model. In Proceedings of the 61st ACM/IEEE Design Automation Conference (DAC), 2024.
- [32] C. Wolf and J. Glaser. Yosys – A free Verilog synthesis suite. In Proceedings of the 21st Austrian Workshop on Microelectronics (Austrochip), 2013.
- [33] K. Xu, J. Sun, Y. Hu, X. Fang, W. Shan, X. Wang, and Z. Jiang. MEIC: Re-thinking RTL debug automation using LLMs. In Proceedings of the 43rd IEEE/ACM International Conference on Computer-Aided Design (ICCAD), 2024. arXiv:2405.06840.
- [34] A. Yang, A. Li, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, et al. Qwen3 technical report, 2025. Qwen Team, Alibaba Group.
- [35] J. Yang, C. E. Jimenez, A. Wettig, K. Lieret, S. Yao, K. Narasimhan, and O. Press.

SWE-agent: Agent-computer interfaces enable automated software engineering. In Advances in Neural Information Processing Systems 37 (NeurIPS), 2024.

- 
- [36] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. ReAct: Synergizing reasoning and acting in language models. In The Eleventh International Conference on Learning Representations (ICLR), 2023. arXiv:2210.03629.
- [37] F. Zhang, B. Chen, Y. Zhang, J. Keung, J. Liu, D. Zan, Y. Mao, J.-G. Lou, and W. Chen. RepoCoder: Repository-level code completion through iterative retrieval and generation. In Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP), pages 2471–2484. Association for Computational Linguistics, 2023.
- [38] Y. Zhao, D. Huang, C. Li, P. Jin, Z. Nan, T. Ma, L. Qi, Y. Pan, Z. Zhang, R. Zhang, X. Zhang, Z. Du, Q. Guo, X. Hu, and Y. Chen. CodeV: Empowering LLMs for Verilog generation through multi-level summarization, 2024.



Appendix A — Prompt Templates

This appendix reproduces the load-bearing prompt text of the system, so that the mechanisms described in Chapters 3 and 4 can be audited at the level at which they actually operate. Text is excerpted verbatim from the repository (`agentic_ip_reuse/.../prompts.py` and `ip_reuse_legacy/prompts.py`); ellipses mark elisions.

A.1 Planner System Prompt (Catalog-Grounding Core)

The planner's system prompt walks the model through the eight planning stages and — decisive in practice — mandates the closed catalog vocabulary:

```
CATALOG GROUNDING (mandatory): The reusable-IP catalog for this
task is listed verbatim in the user message under "Reusable IP
catalog". Treat it as the only source of truth for which IPs
exist.

- Every reuse_decisions entry MUST set "selected_ip" to an
  "ip_id" that appears EXACTLY in that list.
- Never invent, rename, or add suffixes to an IP name (e.g. do
  not turn "e203_exu" into "e203_exu_core").
- If no catalog IP fits a module, set "new_rtl_required": true
  and "selected_ip": null for that module instead of guessing a
  name.
- Produce one reuse_decisions entry for every module that has a
```

```

candidate IP in the catalog.
...
Return final output as one JSON object only, with keys:
requirements, modules, reuse_decisions, integration_plan,
verification_plan, debug_plan, unresolved_assumptions.

```




Listing A.1: Planner system prompt, grounding section (excerpt).

The user message then carries the injected catalog digest (one line per provided module: `ip_id`, port summary, one-sentence description) and closes with: “Every reuse_decisions.selected_ip must be one of the ip_id values listed in the catalog above; do not invent names.” The post-hoc grounding pass of Section 3.2 enforces the same contract on the returned plan, so the closed vocabulary holds regardless of the trajectory the planning agent takes through its tools.

A.2 Planner View vs. Generation View of a Condensed Specification

Large specifications are condensed once into a chunk-and-merge manifest (Section 3.2) and then rendered into two different textual views, one per consumer. Both views share the same skeleton, in this order: top module name; system summary; the top module’s public interface; its submodule instances and connections; shared clocks, resets, parameters, constraints, and assumptions; authoritative interface excerpts carried verbatim from the raw specification (with an instruction to trust them over any summarized interface); every module’s interface listed exactly once; and finally per-module behavioral requirements. The two views differ in exactly the dimensions their consumers need:



SHARED SKELETON (both views)

TOP MODULE / SYSTEM SUMMARY

TOP PUBLIC INTERFACE / SUBMODULE INSTANCES

SHARED CLOCKS, RESETS, PARAMETERS, CONSTRAINTS

AUTHORITATIVE INTERFACE EXCERPTS (verbatim, trusted)

MODULE INTERFACES (each module exactly once)

MODULE BEHAVIORAL REQUIREMENTS

PLANNER VIEW (route/reuse decisions)

- + per-module "reuse_query:" line -- a ready-made catalog search query for each module
- + behavioral requirements digested to the first 3 items per module ("... plus N more requirement(s)")

GENERATION VIEW (writing the RTL)

- + full behavioral requirements for every module that must be generated
- + modules named as provided IP are interface-only, tagged "[provided reusable IP - instantiate, do not regenerate]"
- + section order: to-generate modules first (in dependency order), provided IP last -- so head-truncation under a context budget keeps what the generator has to build

Listing A.2: Planner view vs. generation view: the rendered differences.

The planner view is thus a reuse-structure lens — interfaces, reuse queries, and a short behavioral digest — while the generation view is a construction lens — complete requirements for new logic and interface-only contracts for everything provided. This division is also what makes the plan informative for routing (Section 4.3): module characteristics that are ambiguous in the raw specification arrive in the plan already separated

into delegation structure and own behavior.



A.3 Specifications Routed as Uncertain

Tier-0 of the cascade router deferred exactly five of 60 tasks to the plan probe in the full-cascade run (Section 6.7). For each, the LLM feature extractor returned no class-A signal (delegation, memory encapsulation, thin control) and no class-B signal (FSM, CDC, datapath algorithm), at full confidence — the specifications genuinely commit to neither class. Their opening self-descriptions show the common character:


sd_bd “manages the buffer descriptors used for data transmission and reception between the system memory and the SD card ... implements a circular buffer of descriptors” — a descriptor store with pointer bookkeeping: neither a wrapper nor a nameable algorithm.

sd_data_serial_host “the interface towards the physical SD card device Data port” — a serial engine mixing protocol timing, CRC handshake, and bus-facing buffering.

e203_biu “controls the ICB requests to the external memory system ... serves as a crucial interface between the LSU, IFU, various peripheral interfaces, and the main memory” — pure interface arbitration between protocols.

e203_exu_alu_lsugu “implements the AGU for load/store and AMO instructions ... shares most of its datapath with the ALU” — address generation entangled with a shared datapath, neither self-contained nor delegating.

e203_exu_wbck “arbitrates between the write-back requests from the ALU and long-pipeline instructions” — an arbiter: modest own state, no submodules, no algorithm by name.



The shared shape — interface-arbitration modules that mediate between protocols — is a coherent third type of question sitting between the two classes, and it is exactly the population the plan probe is designed to adjudicate: once a plan exists, the presence or absence of delegation structure becomes explicit.


A.4 Generation Prompt Structure

The generator prompt is assembled from five blocks, in order: (1) the (possibly condensed) specification; (2) the source of every reused module, prefixed with the environment note that these modules are provided as files and must be instantiated, never re-declared; (3) environment notes (macros defined by the build, the one not-shipped dependency); (4) the testbench’s DUT instantiation, verbatim, introduced as the binding port contract; and (5) the output-format instruction (one complete Verilog module in a fenced code block).

A.5 Functional Repair Prompt

The functional repair turn (Section 3.4.2) re-frames the task as logic repair on a design whose interface is already correct. Its distinctive elements:

```
The following {target_hdl} design COMPILES and matches the
required port interface, but its outputs MISMATCH the reference
testbench:
```

```
<verbatim simulator mismatch report, e.g.
```

```
"Output 'b' has 23 mismatches. First mismatch at time 115.">
```

```
Fix the LOGIC only. Do not change the module name or the port  
list. Reused modules listed as provided source files must only  
be instantiated, never re-declared in your output.
```

```
<behavioral slice: specification sections mentioning the  
mismatching output identifiers, plus verbatim interface  
excerpts (Section 3.4.3); full spec if no identifier named>
```

Listing A.3: Functional repair prompt, framing (paraphrased structure; mismatch report inserted verbatim).

When the repair loop runs on the direct flow (no plan), the reuse-framing block is dropped and the same skeleton applies — the mechanism by which the router’s direct arm shares the pipeline’s repair machinery (Section 4.3).



Appendix B — Per-Module Results and Reproduction

B.1 Per-Module Pass Counts

Table B.1 lists, for every module task, the number of passing samples over generated samples in five representative runs (labels per Table 5.2). This is the tabular form of the heatmap in Figure 6.10; the table body is generated directly from each run’s records by `scripts/thesis_figures/gen_module_table.py`.

Table B.1: Passing samples / generated samples per module task. “–” = task absent from the run.

Module	DIRECT-T0	PIPE-T0	CACHE-OFF	PIPE-120B	ROUTER-2T
aes_cipher_top	0/1	0/1	0/20	0/10	0/10
aes_inv_cipher_top	0/1	0/1	0/20	0/10	0/10
aes_inv_sbox	0/1	0/1	0/20	0/10	6/10
aes_key_expand_128	0/1	1/1	20/20	10/10	4/10
aes_rcon	1/1	1/1	6/20	5/10	2/10
aes_sbox	0/1	0/1	0/20	0/10	3/10
sd_bd	0/1	0/1	0/20	0/10	0/10
sd_clock_divider	1/1	1/1	0/20	0/10	0/10
sd_cmd_master	0/1	0/1	0/20	0/10	0/10
sd_cmd_serial_host	0/1	0/1	0/20	0/10	0/10


Module	DIRECT-T0	PIPE-T0	CACHE-OFF	PIPE-120B	ROUTER-2T
sd_controller_wb	0/1	0/1	0/20	0/10	0/10
sd_crc_16	1/1	1/1	18/20	8/10	10/10
sd_crc_7	1/1	1/1	18/20	9/10	10/10
sd_data_master	0/1	0/1	0/20	0/10	0/10
sd_data_serial_host	0/1	0/1	0/20	0/10	0/10
sd_fifo_rx_filler	0/1	0/1	1/20	0/10	0/10
sd_fifo_tx_filler	0/1	1/1	19/20	8/10	10/10
sd_rx_fifo	0/1	0/1	0/20	0/10	0/10
sd_tx_fifo	0/1	0/1	0/20	0/10	0/10
sdc_controller	0/1	0/1	0/20	0/10	0/10
e203_biu	0/1	0/1	0/20	0/10	0/10
e203_clk_ctrl	0/1	1/1	9/20	1/10	7/10
e203_clkgate	1/1	1/1	20/20	10/10	10/10
e203_core	0/1	0/1	0/20	0/10	0/10
e203_cpu	0/1	0/1	0/20	0/10	0/10
e203_cpu_top	0/1	0/1	0/20	0/10	0/10
e203_dtcn_ctrl	0/1	0/1	1/20	5/10	4/10
e203_dtcn_ram	0/1	1/1	10/20	10/10	8/10
e203_extend_csr	0/1	1/1	15/20	9/10	6/10
e203_exu	0/1	0/1	0/20	0/10	0/10
e203_exu_alu	0/1	0/1	0/20	0/10	0/10
e203_exu_alu_bjp	0/1	0/1	2/20	2/10	0/10
e203_exu_alu_csrctrl	0/1	0/1	0/20	4/10	2/10
e203_exu_alu_dpath	0/1	0/1	0/20	0/10	0/10
e203_exu_alu_lsugu	0/1	0/1	0/20	0/10	0/10
e203_exu_alu_muldiv	0/1	0/1	0/20	0/10	0/10
e203_exu_alu_rglr	0/1	0/1	5/20	7/10	5/10
e203_exu_branchslv	0/1	0/1	0/20	0/10	0/10
e203_exu_commit	0/1	0/1	0/20	1/10	0/10
e203_exu_csr	0/1	0/1	0/20	0/10	0/10

Module	DIRECT-T0	PIPE-T0	CACHE-OFF	PIPE-120B	ROUTER-2T
e203_exu_decode	0/1	0/1	0/20	0/10	0/10
e203_exu_disp	0/1	0/1	0/20	0/10	1/10
e203_exu_excp	0/1	0/1	0/20	0/10	0/10
e203_exu_longpwbck	0/1	0/1	0/20	0/10	0/10
e203_exu_nice	0/1	0/1	0/20	0/10	0/10
e203_exu_oitf	0/1	0/1	0/20	3/10	0/10
e203_exu_regfile	0/1	1/1	2/20	5/10	6/10
e203_exu_wbck	0/1	0/1	0/20	0/10	0/10
e203_ifu	0/1	0/1	14/20	10/10	5/10
e203_ifu_ifetch	0/1	0/1	0/20	0/10	0/10
e203_ifu_ift2icb	0/1	0/1	0/20	0/10	0/10
e203_ifu_litebpu	0/1	0/1	0/20	0/10	0/10
e203_ifu_minidec	0/1	1/1	20/20	10/10	9/10
e203_irq_sync	0/1	1/1	14/20	10/10	7/10
e203_itcm_ctrl	0/1	0/1	0/20	0/10	0/10
e203_itcm_ram	0/1	1/1	8/20	10/10	5/10
e203_lsu	0/1	0/1	0/20	0/10	0/10
e203_lsu_ctrl	0/1	0/1	0/20	0/10	0/10
e203_reset_ctrl	0/1	0/1	0/20	0/10	0/10
e203_srams	0/1	1/1	20/20	10/10	10/10

B.2 Reproduction Commands

All runners live under scripts/ in the project repository and require a live vLLM endpoint (VLLM_BASE_URL, VLLM_MODEL) plus Verilator on PATH. The commands below reproduce the thesis’s principal arms; output directories match Table 5.3.

```
# Headline pipeline (temp 0) and pure-model baseline
```



```
python scripts/run_agentic_plan_legacy_realbench.py \
  --task-level module --samples 1 --temperature 0 \
  --output-dir runs/agentic_plan_legacy_realbench_plan_hallu_fix_t0
python scripts/run_realbench_direct_model.py \
  --task-level module --samples 1 --temperature 0 \
  --output-dir runs/realbench_direct_model

# Repair-cache ablation (swap off | task | run)
python scripts/run_agentic_plan_legacy_realbench.py \
  --task-level module --samples 20 \
  --planner-search-mode hybrid --repair-cache task \
  --output-dir runs/..._task_rep_cache

# Grounding single-fix ablations
# --no-planner-inject-catalog | --no-planner-ground-reuse |
# --no-planner-completeness-gate

# Full cascade router (Tier-0 LLM decider + Tier-1 plan probe)
python scripts/run_realbench_routed.py \
  --router cascade --decider llm --samples 10 \
  --output-dir runs/Full-2T_router_oss20B_sync6_func4

# Per-module pass tables and figures
python scripts/module_pass_rate_report.py runs/<run_dir> \
  -o <run_dir>/module_pass_rate_analysis.md
cd scripts/thesis_figures && for f in fig_*.py; do python $f; done
```

Listing B.4: Reproducing the principal experimental arms.